ParlayANN: Scalable and Deterministic Parallel Graph-Based Approximate Nearest Neighbor Search Algorithms

Magdalen Dobson Manohar, Zheqi shen, Guy E. Blelloch, Laxman Dhulipala, Yan Gu, Harsha Vardhan Simhadri, Yihan Sun PPoPP 2024

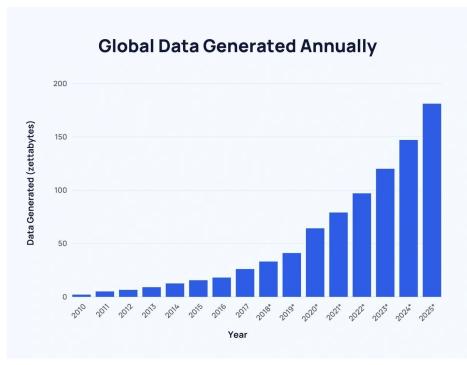
> 2024. 10. 23 Hojin Shin (Dankook University) ghwls03s@gmail.com or hojin03s@dankook.ac.kr





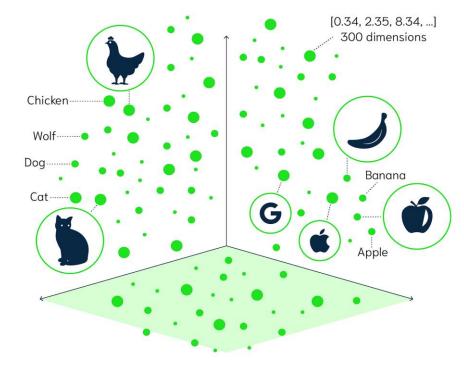
Contents

- 1. Introduction
- 2. Background
- 3. Motivation
- 4. ParlayANN (Libaray)
- 5. Evaluation
- 6. Conclusion



source:https://www.statista.com/statistics/871513/worldwide-data-created/

Constantly growing amount of the available information resources



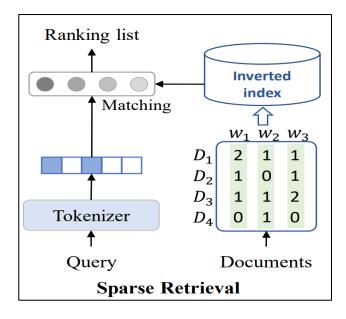
source:https://medium.com/@serkan_ozal/vector-similarity-search-53ed42b951d9



High demand in scalable and efficient similarity search data structures

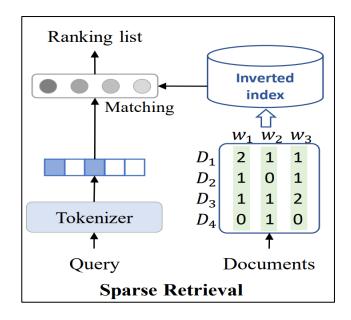






Consider *only some words*in the target documents and search terms among a large number of words



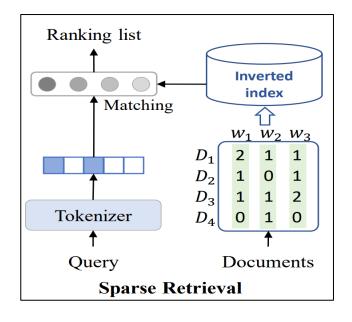


Consider *only some words* in the target documents and search terms among a large number of words

have similar meanings
but do not match

경기도민의 수는 얼마?

경기도의 인구는 00명

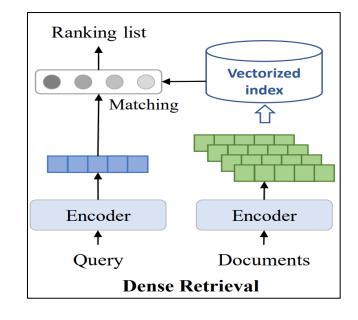


Consider *only some words* in the target documents and search terms among a large number of words

have similar meanings
but do not match

경기도민의 수는 얼마?

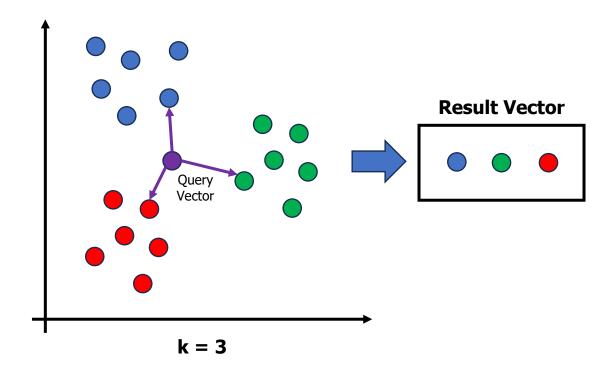
경기도의 인구는 00명



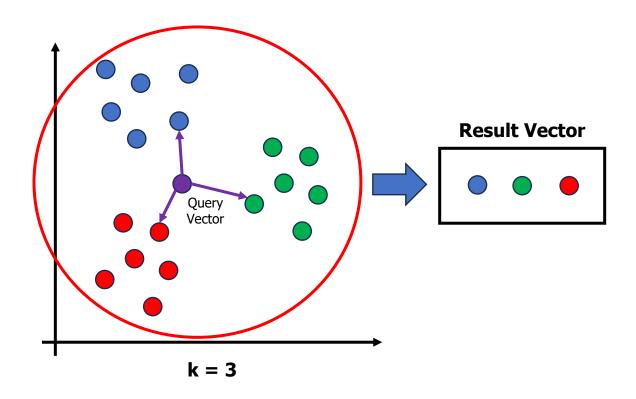
Continuous **vector spaces**, capturing **semantic similarity** beyond exact term matching



Similarity Search (K-NNS)



Similarity Search (K-NNS)



Strengths

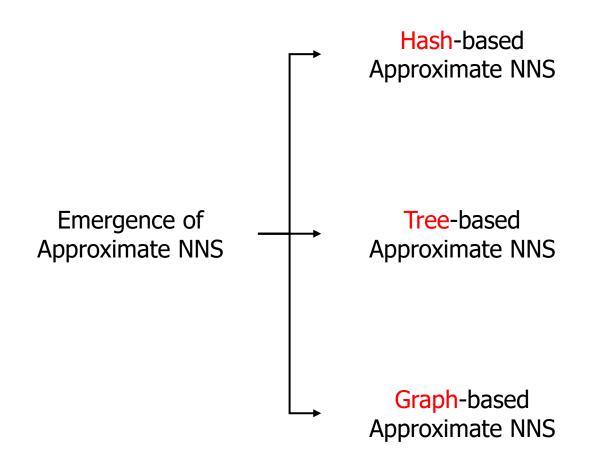
Exact nearest neighbor search→ High *Accuracy*

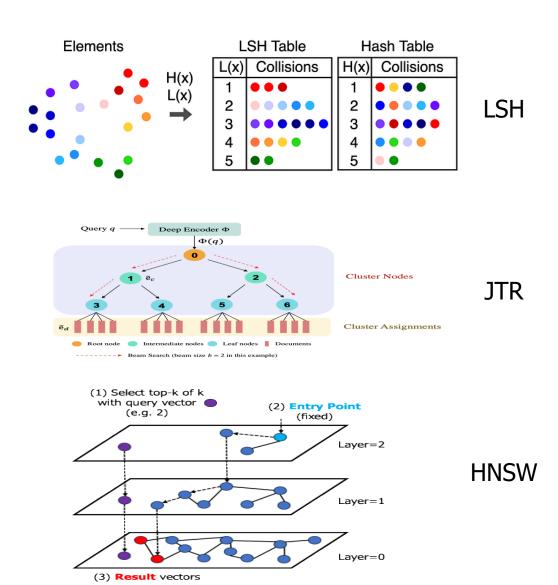
Limitations

The number of stored elements making it *infeasible* for large-scale dataset (*High-D*)

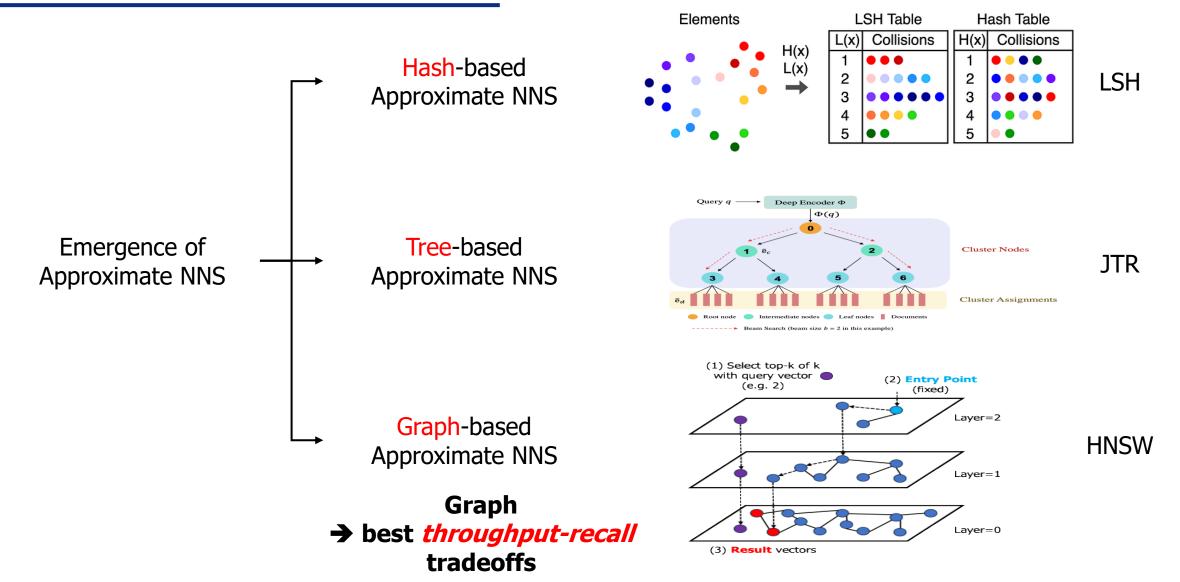
→ Whole search







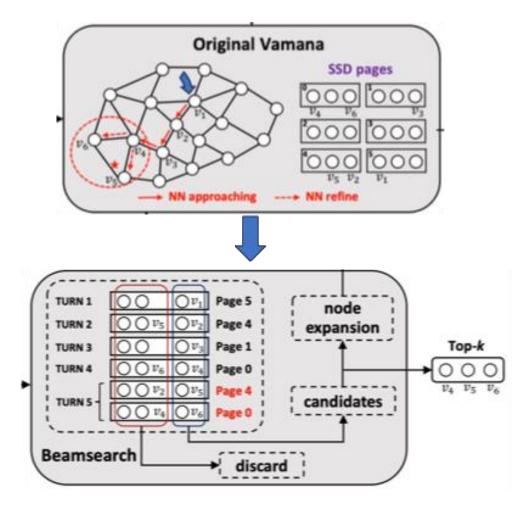




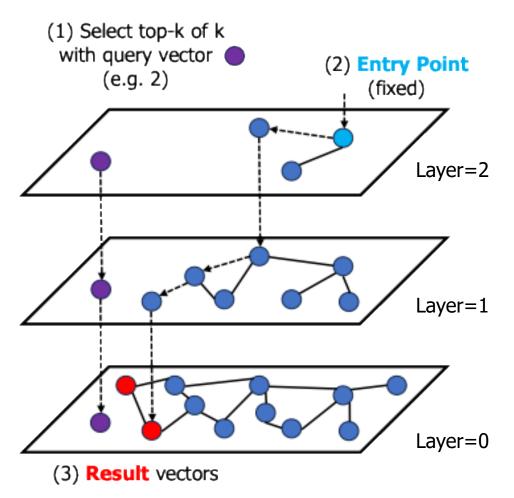




2. Background



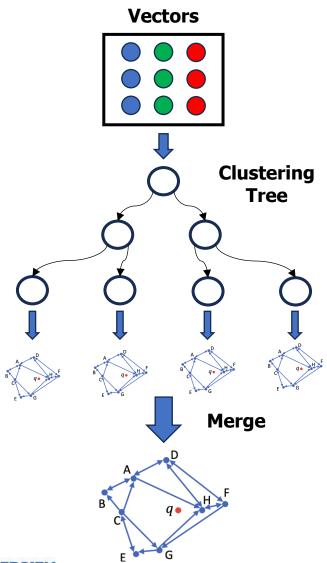
DiskANN (Vamana)



HNSW Hierarchical Navigable Small World



2. Background



HCNNG (Hierarchical Clustering-based Nearest Neighbor Graph)

Hierarchical Clustering-Based Graphs for Large Scale Approximate Nearest Neighbor Search

Javier Vargas Muñoz a,*, Marcos A. Gonçalves b, Zanoni Dias a, Ricardo da S. Torres a

PyNNDescent (Python Nearest Neigbhor Descent)

Efficient K-Nearest Neighbor Graph Construction for Generic Similarity Measures

Wei Dong Moses Charikar wdong@cs.princeton.edu moses@cs.princeton.edu

Kai Li li@cs.princeton.edu

Department of Computer Science, Princeton University 35 Olden Street, Princeton, NJ 08540, USA





a University of Campinas, Av. Albert Einstein, 1251, Campinas, São Paulo, Brazil

^b Universidade Federal de Minas Gerais, Av. Antônio Carlos, 6627, Belo Horizonte, Minas Gerais, Brazil

3. Motivation

Performance Issues on Large Datasets

- High-dimensional vector embeddings
 Essential element
- Transform into millions or billions of vectors
- Graph-based ANNS suffer from scalability issues on large datasets
- During parallel processing due to use of locks (per-point locks)

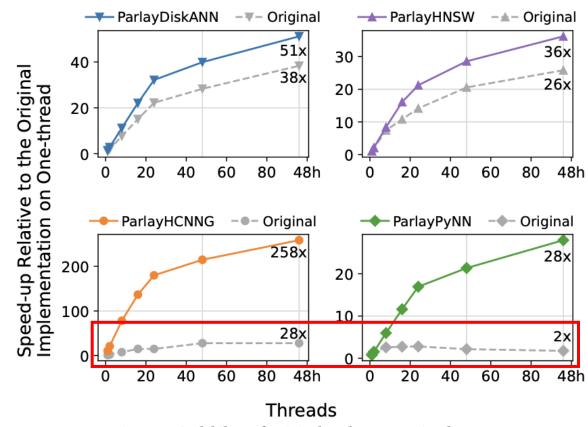


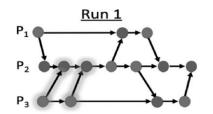
Figure 1. Scalability of original and our new implementations of four ANNS algorithms on various number of threads. Within each subfigure, all numbers are speedup numbers relative to the original implementation on one thread. Higher is better. Results were tested on a machine with 48 cores using dataset BIGANN-1M (10⁶ points). "48h": 48 cores with hyperthreads. The two implementations in the same subfigure always use the same parameters and give similar query quality (recall-QPS curve).

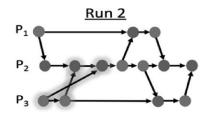


3. Motivation

Use of Locks Limitations

- Existing parallel ANNS use locks to prevent data conflicts
- Locks are effective in maintaining data consistency
- The execution results depending on the order in which the locks are acquired
- Application (persistence, crash recovery)
- Non-determinism





Existing Benchmarks Limitations

 Most of the existing studies evaluate relatively *small dataset* of millions of data

A Comprehensive Survey and Experimental Comparison of Graph-Based Approximate Nearest Neighbor Search

Mengzhao Wang¹, Xiaoliang Xu¹, Qiang Yue¹, Yuxiang Wang^{1,*}

¹Hangzhou Dianzi University, China

{mzwang,xxl,yq,lsswyx}@hdu.edu.cn

ANN-Benchmarks: A Benchmarking Tool for Approximate Nearest Neighbor Algorithms*

Martin Aumüller¹, Erik Bernhardsson², and Alexander Faithfull¹

¹ IT University of Copenhagen, Denmark, {maau,alef}@itu.dk
² Better, mail@erikbern.com

Table 3: Statistics of real-world datasets.

| Dataset | Dimension | # Base | # Query | LID [34, 57] |
|------------|-----------|-----------|---------|--------------|
| UQ-V [5] | 256 | 1,000,000 | 10,000 | 7.2 |
| Msong [2] | 420 | 992,272 | 200 | 9.5 |
| Audio [4] | 192 | 53,387 | 200 | 5.6 |
| SIFT1M [1] | 128 | 1,000,000 | 10,000 | 9.3 |
| GIST1M [1] | 960 | 1,000,000 | 1,000 | 18.9 |
| Crawl [3] | 300 | 1,989,995 | 10,000 | 15.7 |
| GloVe [48] | 100 | 1,183,514 | 10,000 | 20.0 |
| Enron [75] | 1,369 | 94,987 | 200 | 11.7 |

| Dataset | Data/Query Points | Dimensionality | Metric |
|-----------------|--------------------|----------------|-----------|
| SIFT | 1000000 / 10000 | 128 | Euclidean |
| GLOVE | 1 183 514 / 10 000 | 100 | Angular |
| NYTimes | 234791 / 10000 | 256 | Euclidean |
| Rand-Angular | 1000000 / 1000 | 128 | Angular |
| SIFT-Hamming | 1000000/10000 | 256 | Hamming |
| NYTimes-Hamming | $234791\ /\ 10000$ | 128 | Hamming |

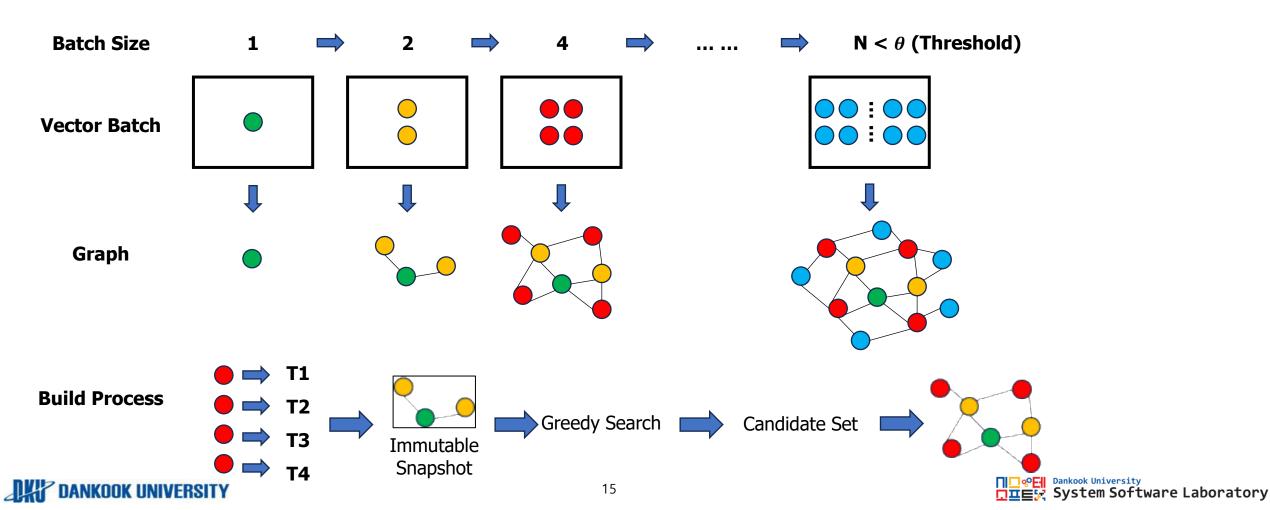
Table 3. Datasets under consideration





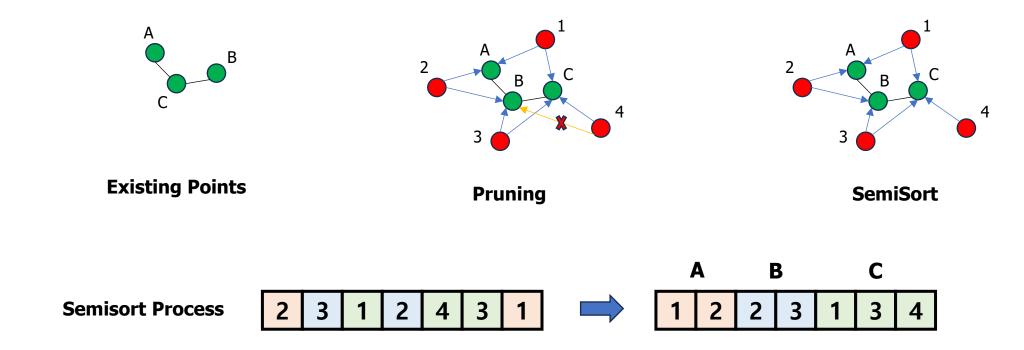
(1) Batch Insertion

Divide the data points into batches and insert them incrementally/independently



(2) Pruning and SemiSort

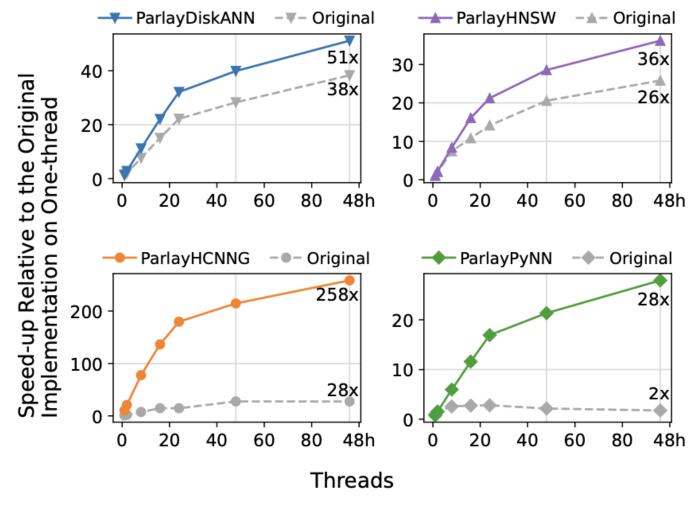
- Remove unnecessary connections and reduce the complexity of the graph
- Groups items by key, arranging with the same key in a sequence



4. ParlayANN

```
Algorithm 3: batchBuild(\mathcal{P}, s, R, L).
    Input: Point set \mathcal{P}, starting point s, beam width L, degree
              bound R.
    Output: An ANN graph consisting of all points in \mathcal{P}.
1 start \leftarrow 1
2 while start \leq |\mathcal{P}| do
                                                                     // Prefix-doubling
          end \leftarrow \min(start \times 2, start + \theta, |\mathcal{P}|) // \theta: batch size upper
            bound
          BatchInsert(\mathcal{P}[start..end])
                                                                    Batch Insertion
          start \leftarrow end + 1
6 Function BatchInsert(\mathcal{P}') // Insert a batch \mathcal{P}' to the current
      index
          parallel for p \in \mathcal{P}' do
 7
                \mathcal{V}.\mathcal{K} \leftarrow \text{greedySearch}(p, s, L, 1)
                N_{\text{out}}(p) \leftarrow \text{prune}(p, \mathcal{V}, R) Pruning
          \mathcal{B} \leftarrow \bigcup_{p \in \mathcal{P}'} N_{\text{out}}(p) // All (existing) affected points
10
          parallel for b \in \mathcal{B} do
11
                // N: all points in P' that added b as their neighbors
                \mathcal{N} \leftarrow \{p \mid p \in \mathcal{P}' \land b \in N_{\text{out}}(p)\}
12
                N_{\text{out}}(b) \leftarrow N_{\text{out}}(b) \cup \mathcal{N}
                                                                           SemiSort
13
                if |N_{out}(b)| > R then N_{out}(b) \leftarrow
14
                  prune(b, N_{out}(b), R)
```

Comparison with multi-thread





Comparison with another library

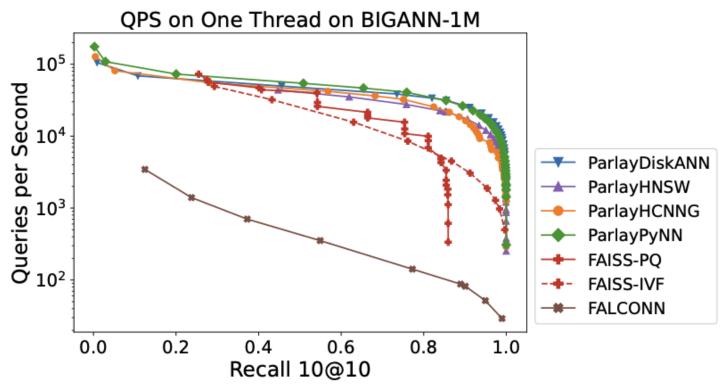


Figure 5. QPS on a single thread on BIGANN-1M. Shown to compare with ANN-benchmarks.



Comparison with Billion-scale Dataset

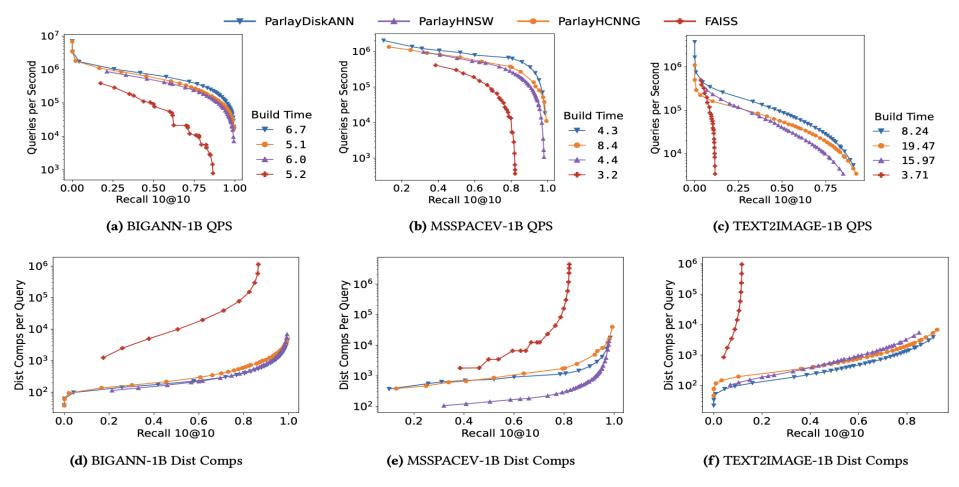


Figure 3. Build time (hours), QPS, recall, and distance comparisons for all algorithms on billion-size datasets.



Comparison with Million-scale Dataset

BIGANN MSSPACEV TEXT2IMAGE DiskANN .42 .35 .70 **HNSW** .35 .37 .94 **HCNNG** .45 .77 1.75 pyNNDescent .42 .73 1.23 **FAISS** .19 .13 .22

Table 1. Build times (hours) on hundred million scale datasets.

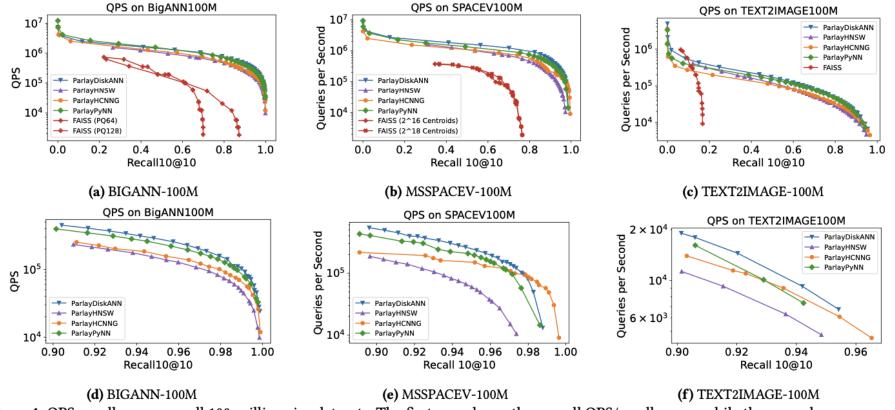


Figure 4. QPS-recall curves on all 100-million size datasets. The first row shows the overall QPS/recall curve, while the second row zooms into a higher-recall regime. The build times are given in Tab. 1



6. Conclusion

Present **ParlayANN**, which implements four **parallel deterministic** graph-based ANNS algorithms

ParalyANN that scale to **billion-scale inputs** on a single machine with high recall

Avoid the use of locks, achieve better scalability than existing implementations

Outperformed existing **non-graph** implementations in the ability of achieving high recall



ParlayANN: Scalable and Deterministic Parallel Graph-Based Approximate Nearest Neighbor Search Algorithms

Magdalen Dobson Manohar, Zheqi shen, Guy E. Blelloch, Laxman Dhulipala, Yan Gu, Harsha Vardhan Simhadri, Yihan Sun PPoPP 2024

Thank you!

2024. 10. 23 Hojin Shin (Dankook University) ghwls03s@gmail.com or hojin03s@dankook.ac.kr



