



LearnedKV: Integrating LSM and Learned Index for Superior Performance on SSD

- Wang, Wenlong, and David Hung-Chang Du.
- University of Minnesota

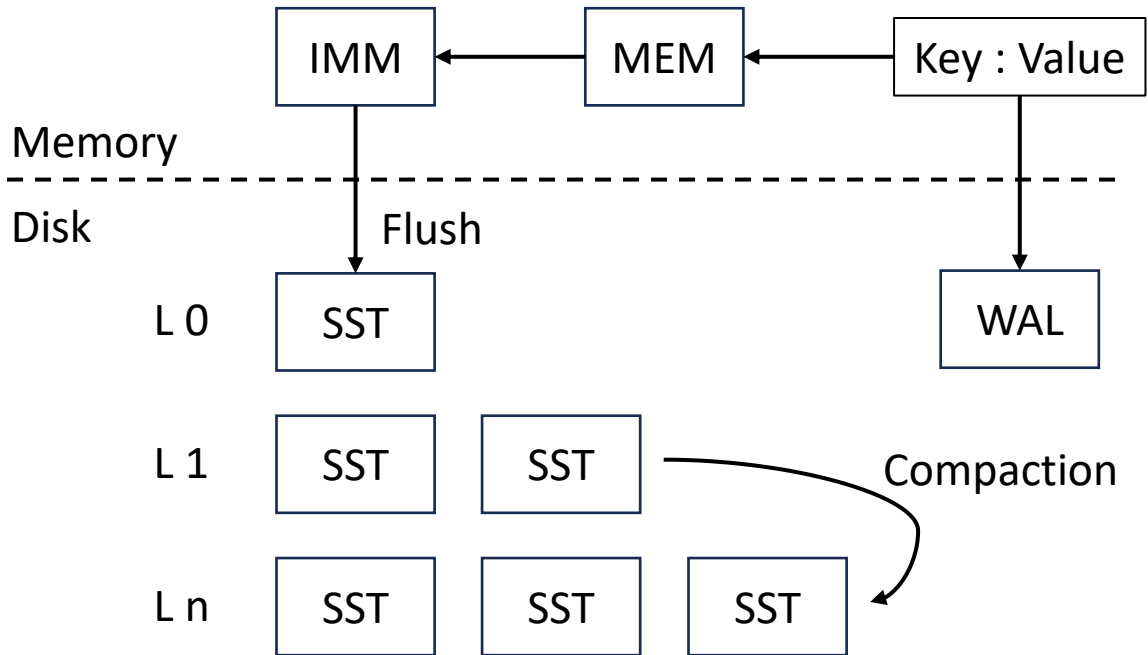
2025.1.22

Presentation by Guangxun Zhao
GuangxunZhao@dankook.ac.kr

Introduction



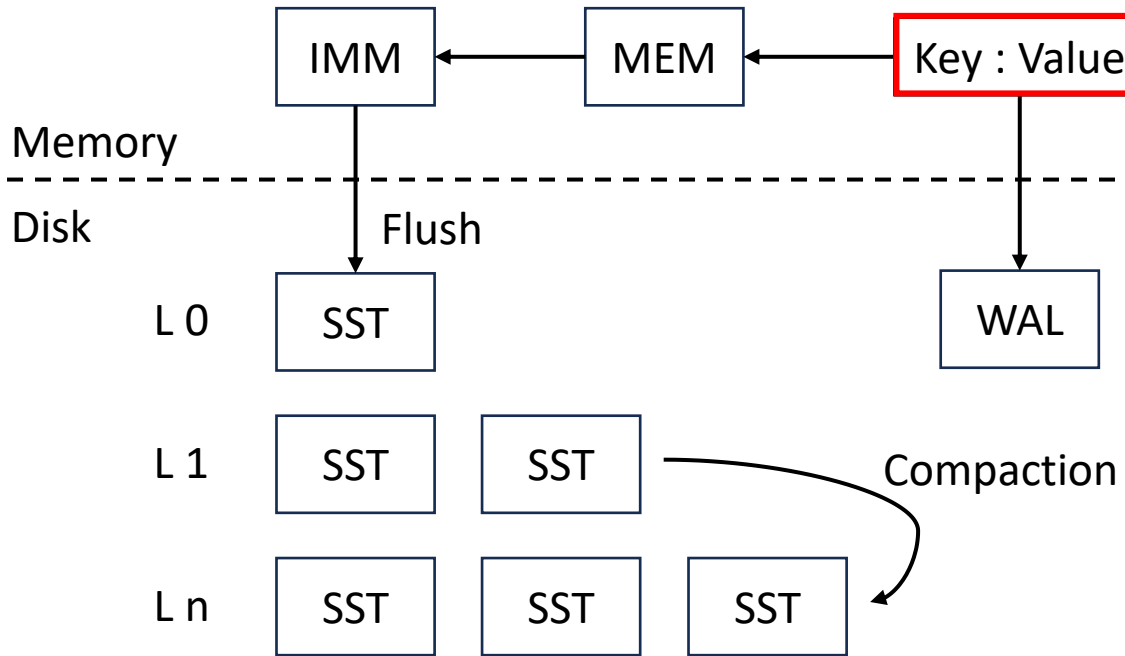
- To effectively handle bigdata use Key-Value Stores.
- LSM-tree is write optimized data structure.



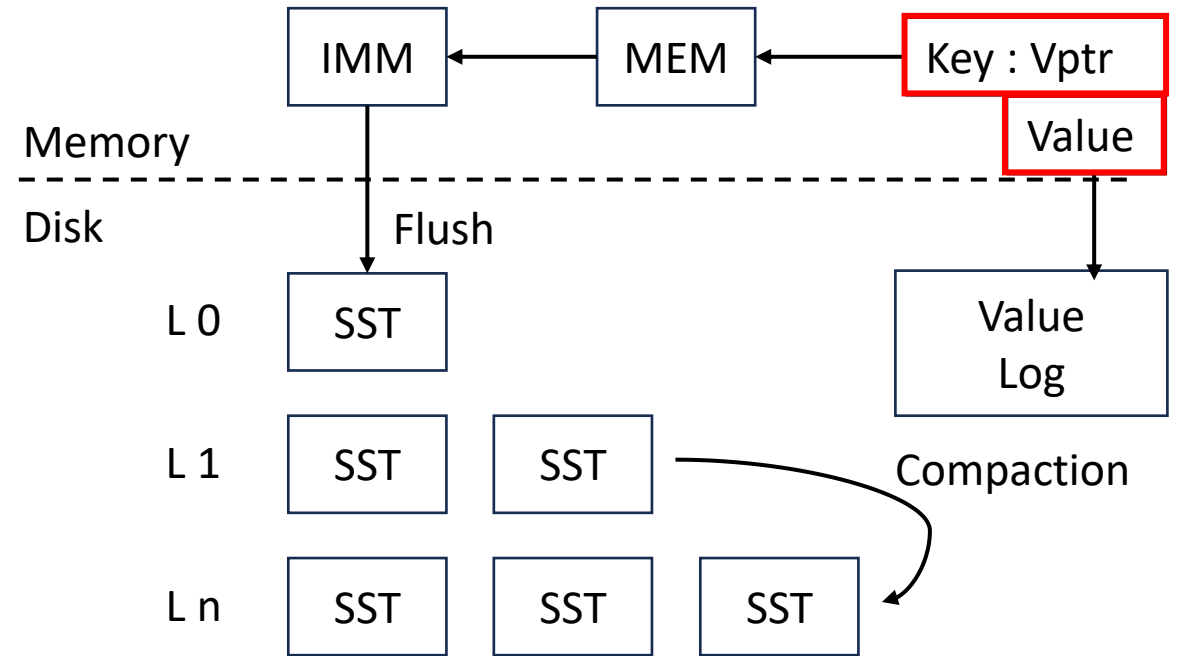
RocksDB and Wisckey Structure



- Write/Read amplification problem in LSM-tree.
- Wisckey uses Key-Value separation to reduce Write/Read amplification .

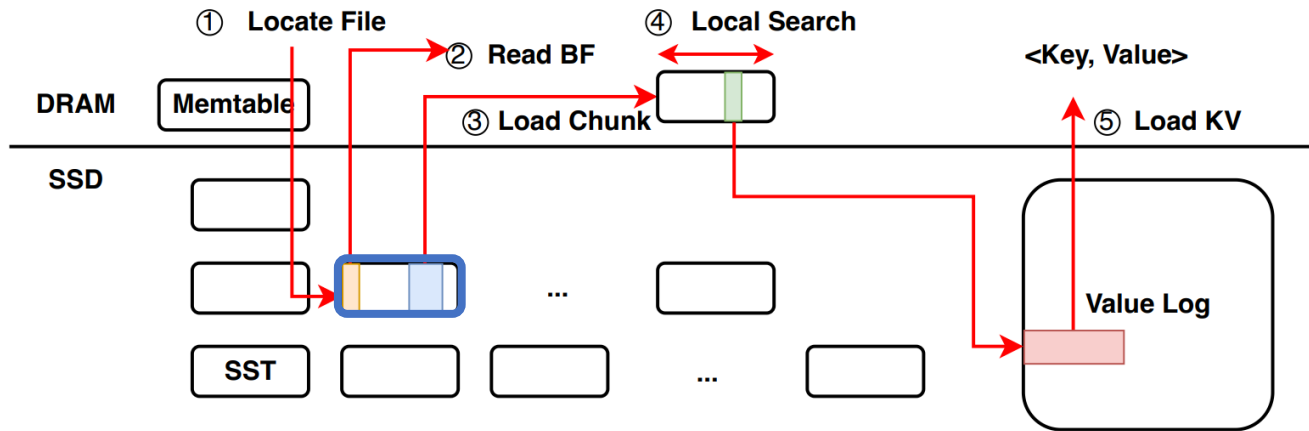


Structure of RocksDB



Structure of Wisckey

- Learning Algorithm: Greedy-PLR.
- Bourbon using the Learned Index to replace or bypass the index blocks of SST files.



Bourbon lookup process

Piecewise Linear Regression

$$\text{Pos} = \text{slop} * \text{key} + \text{intercept}$$

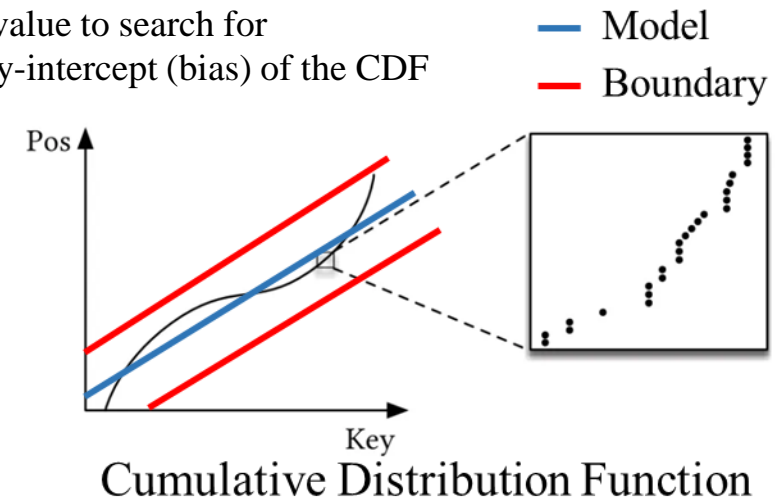
Pos - err , Pos + err

Pos: Position where the key is stored (index in the predicted array)

slope: Slope of the CDF

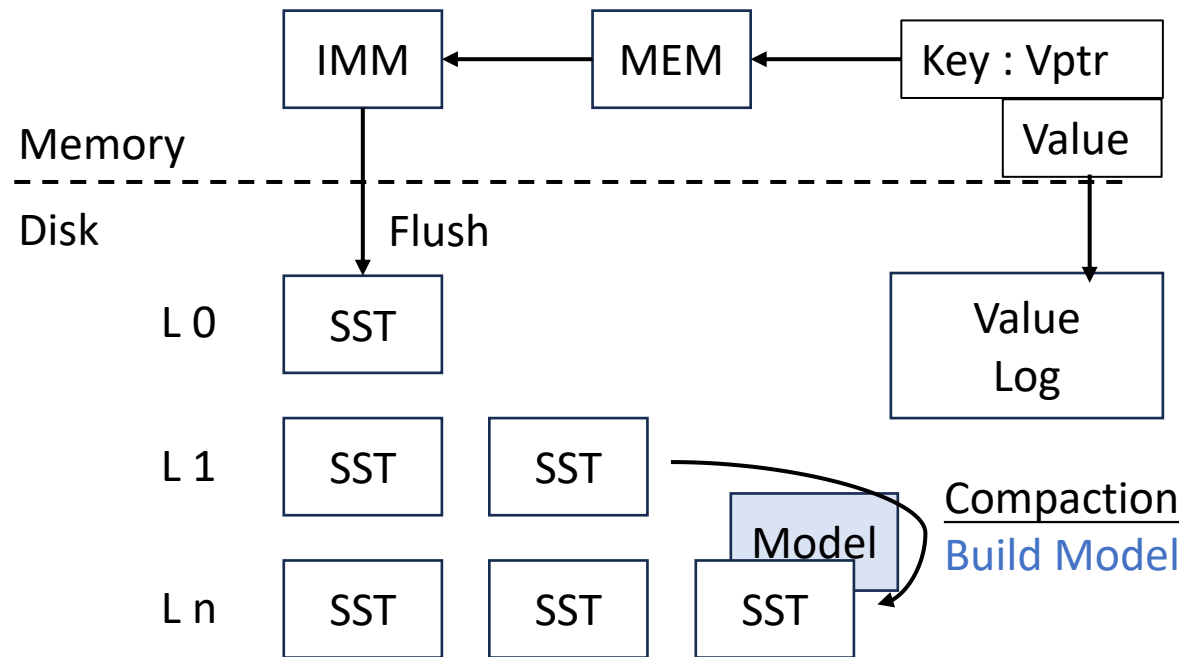
key: Key value to search for

intercept: y-intercept (bias) of the CDF





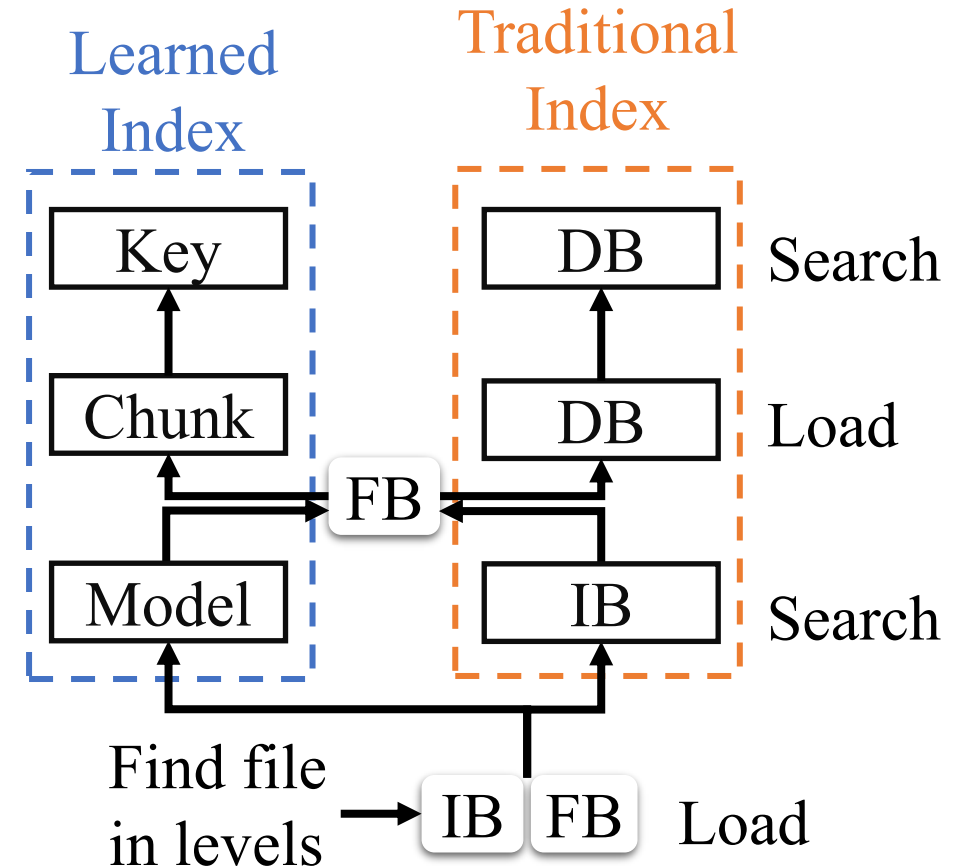

- Models need to be rebuilt during compaction.
- The design merely adds the Learned Index instead of treating it as an independent replacement to optimize the system.



Read and Write time cost of the RocksDB



- 1M workload into 10 stages in RocksDB, writing 100K new keys into it for each stage.



Tiered index



- Aim to build a tiered index where the LSM-tree is used for absorbing random writes, and the Learned Index helps accelerate the lookup process.
 1. Lookup performance
 2. Compaction cost
 3. Re-insert after Garbage Collection
 4. Update latency

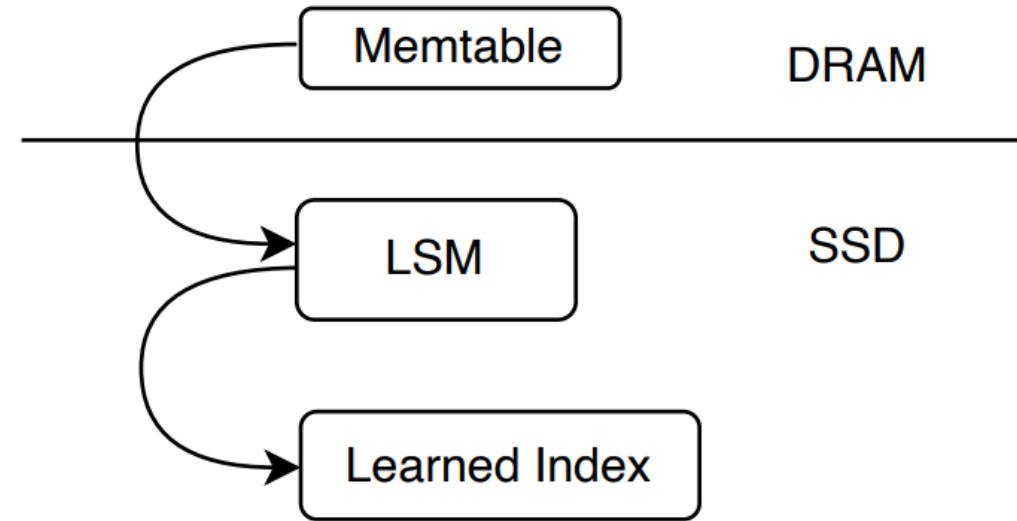


Figure 4: Abstraction of tiered storage with LSM and Learned Index

Challenges



- Challenges of integrate LSM with learned indexes
 - Challenge 1: How to efficiently convert data from the LSM-tree to the Learned Index?
 - Challenge 2: Under which conditions does the conversion provide more benefits?

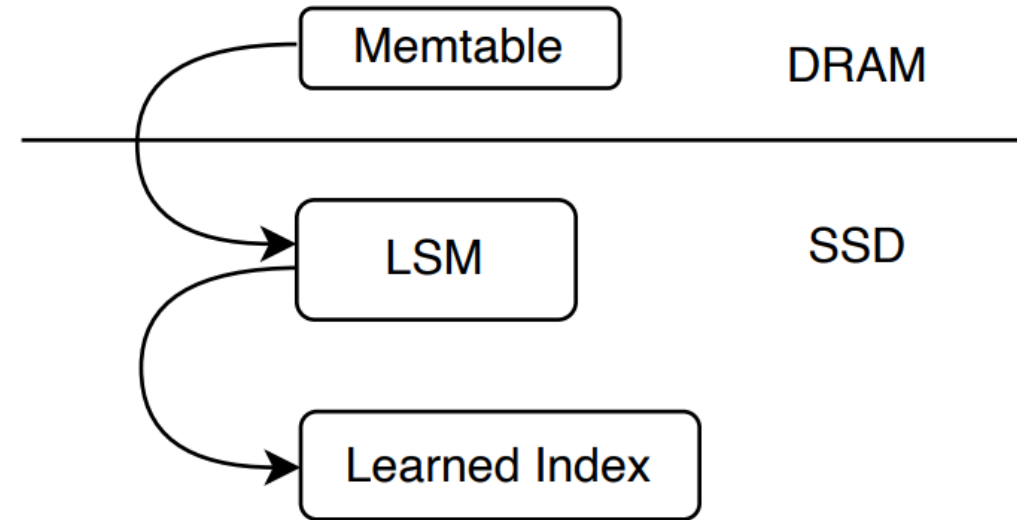
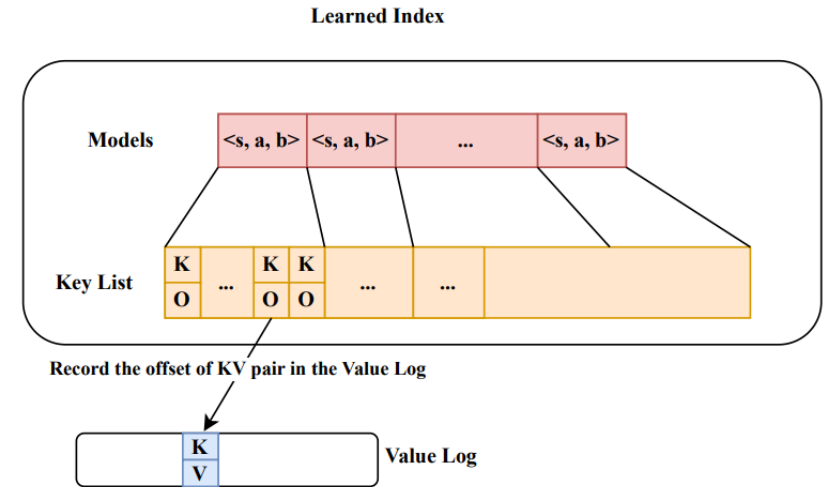
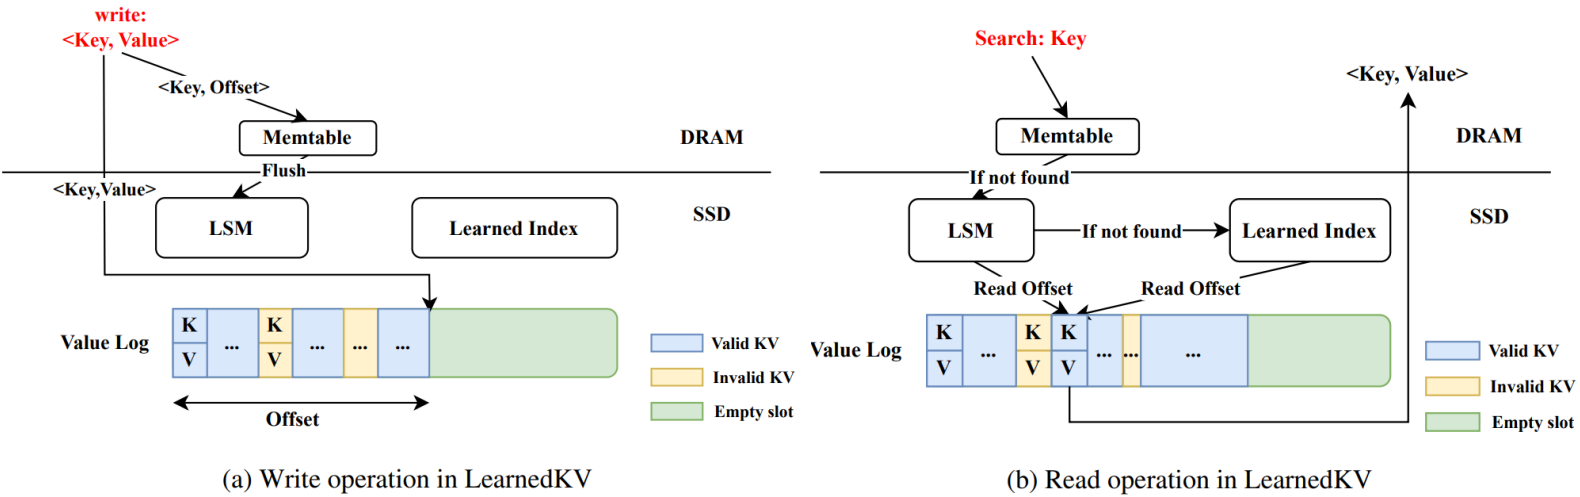


Figure 4: Abstraction of tiered storage with LSM and Learned Index

LearnedKV Structure



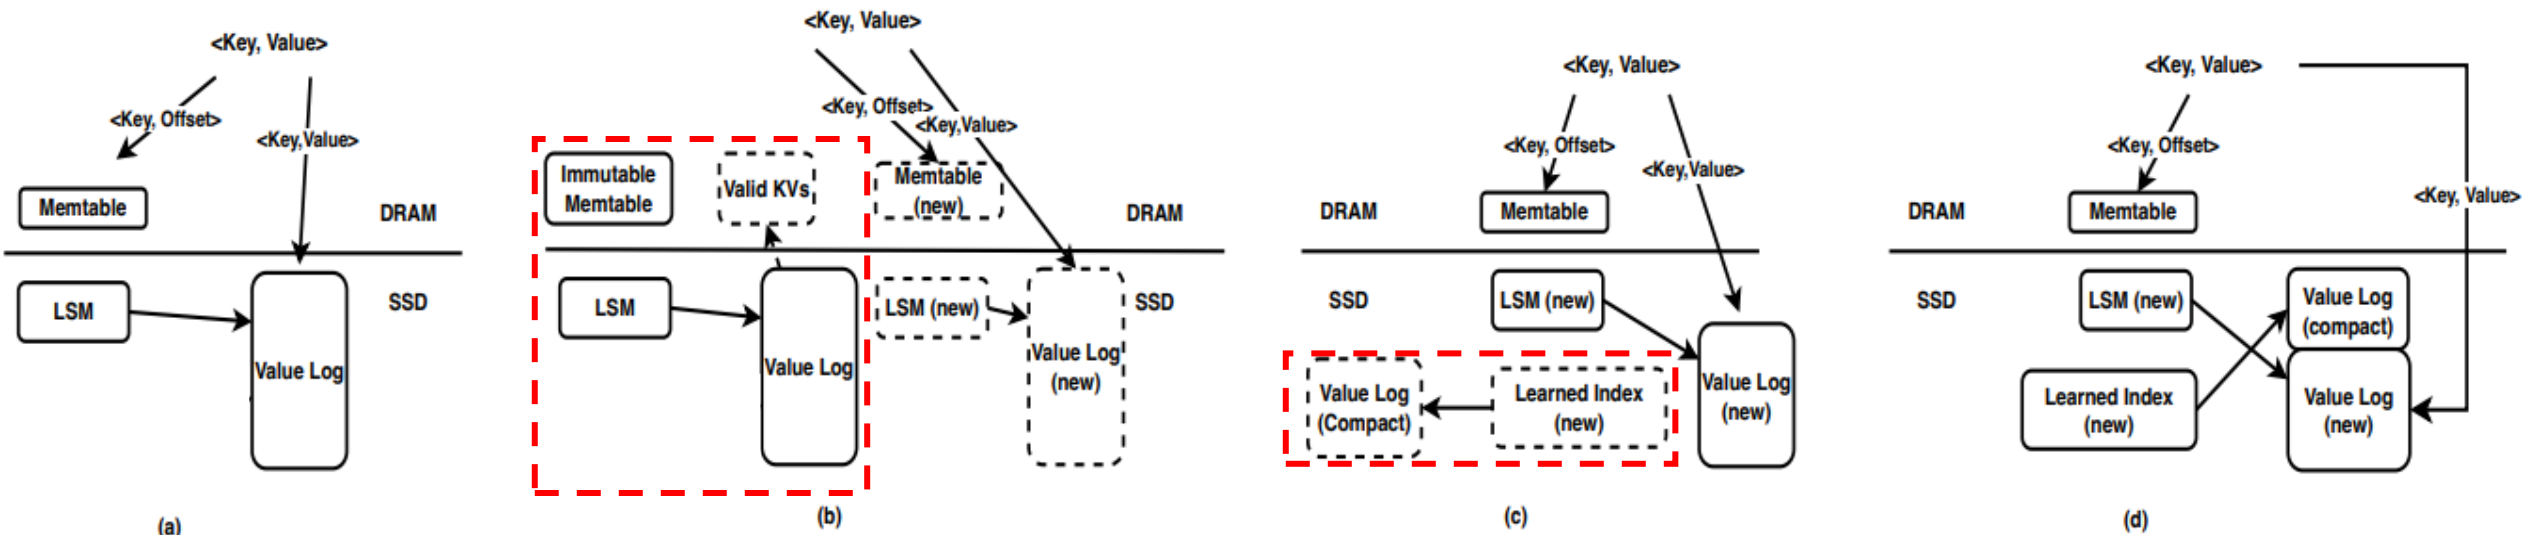
- LearnedKV consists of three main components:
 - LSM-Tree, Learned Index, VLog.



LearnedKV Structure



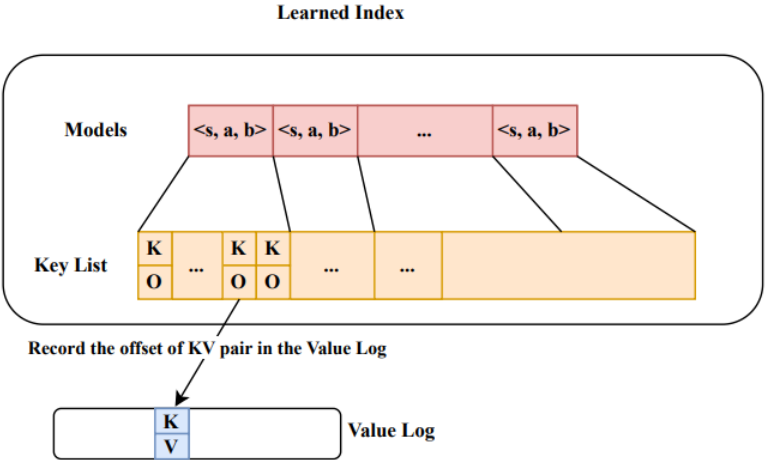
- GC-triggered Conversion from LSM-tree to Learned Index.
- When performing Vlog GC, sorting and learning are applied. After that, the old LSM-tree is deleted and replaced with a new LSM-tree.



Greedy-PLR+



- Learned index
 - Starting key (s), slope (a), and intercept (b).
- $Pos = a \times k + b$
 - Pos=8000
 - Error bound=1
 - Page number : $0,8000 // 4096 = 1, 2$

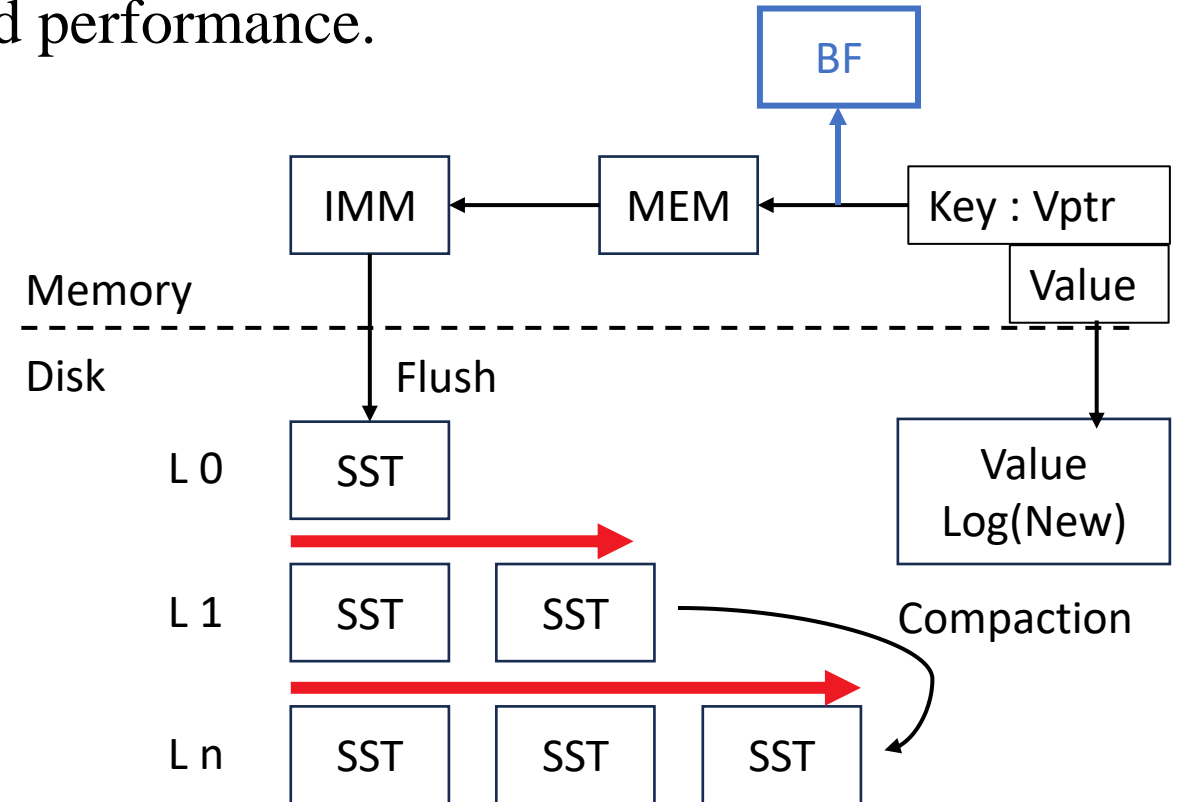
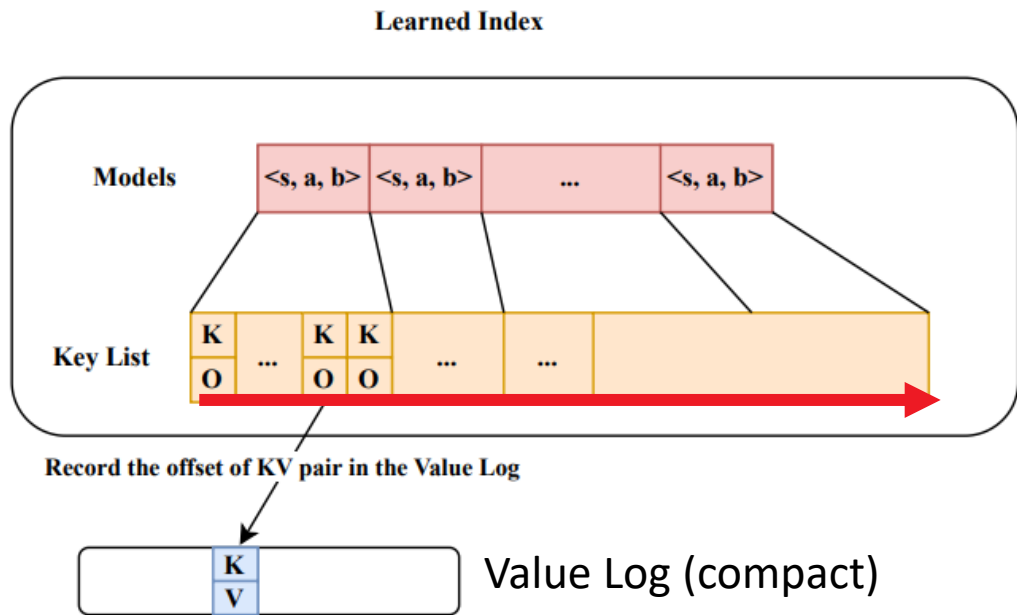


	Bourbon	LearnedKV
예측 모델 방식	바이트 오프셋을 직접 예측	페이지 단위로 예측
검색 범위	[p-e,p+e]	[p-e,p+e] (SSD 페이지 단위)
SSD 접근 방식	오프셋 기반 페이지 로드	페이지 단위에서 직접 예측 후 로드

LearnedKV structure



- Range query optimize with Learned index's Key list.
- Use In-memory Bloom filter to optimize read performance.





Component	Specification	Schemes Compare
Processor	12- core Intel Xeon E5-2620 v3 2.40GHz CPU	BlobDB within RocksDB
Memory	62GB	HashKV(KV separation DB)
Storage	854GB SSD	B+-Tree
key & value size	8 bytes & 1016 bytes	

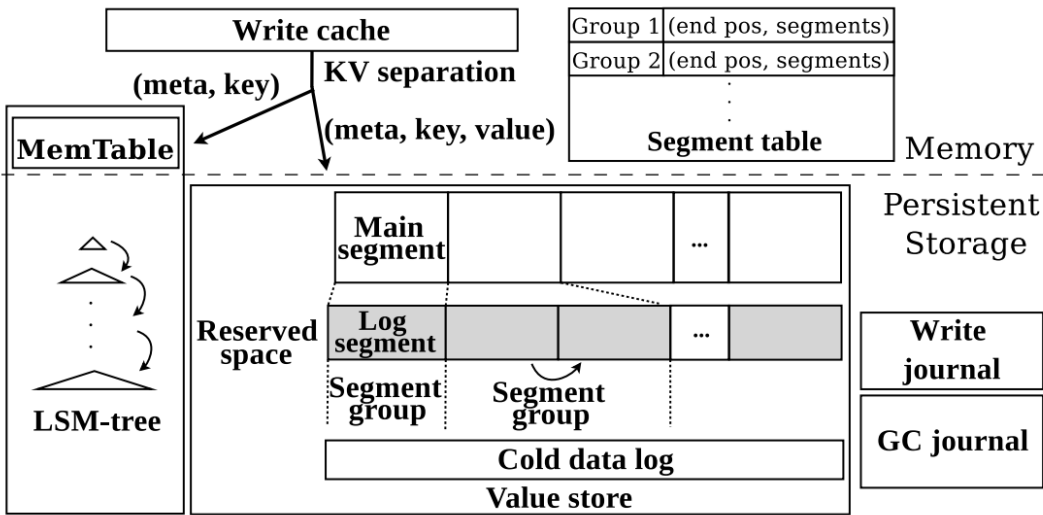


Figure 3: HashKV architecture.

Overall Performance Comparison



- P0 load 1M KV pairs.(YCSB load)
- P1,P2,P3 500,000 read and update operations. (YCSB A)
 - Zipfian distribution.




Figure 8: Throughput Comparison between LearnedKV and RocksDB. (P0: Load; P1,P2,P3: Read and Update; Learned Index is built in the middle of P1; GC happens in P1,P2,P3)

	LearnedKV	RocksDB
Key-Ptr Space		
LSM	4.2MB	23MB
key_array	7.7MB	-
model	8.0KB	-
} Learned index		
Total	11.9MB	23MB
Key-Value Space		
vlog	978MB	978MB

Overall Performance Comparison



- Each scan request reads about 500KB of KV pairs.
- "Set 1" contains 0.5M updates before the scan requests, while "Set 2" includes 1M update requests before the scan.




Figure 10: Performance Comparison of Range Scan

Other KV Stores



- LearnedKV can also outperform HashKV 3.59x in terms of write performance
- With an 8-byte value setup, LearnedKV outperforms the B+-Tree by more than 5 times in both read and write performance.




Figure 11: Comparison between LearnedKV and HashKV




Figure 12: Comparison between LearnedKV and B+_Tree

Drill-down Analysis



- Read and write operation numbers and their average time cost in the P3.
- 500K read operations in Read time break down.

Operation	Number	Avg Time Cost (μs)
Write	500,377	165.584
Read	499,623	66.6561
Read Operation		
Read through LI	131,508 26	37.5138
Read through RD	383,256 %	58.0299
Load form vlog	499,623	11.6553

Table 2: Summary of operations and their time costs. (LI: Learned Index; RD: RocksDB)


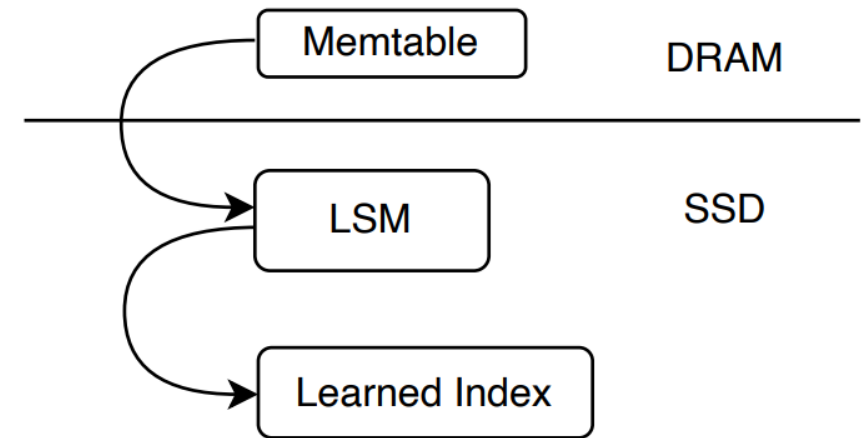


Figure 13: Comparison of Throughput With and Without Bloom Filter

Conclusion



- This paper proposes LearnedKV, an efficient tiered KV store that integrates an LSM with a Learned Index.
- LearnedKV standalone design significantly reduces the size of the LSM, further enhancing both read and write performance.
- LearnedKV outperforms the latest KV systems, achieving up to 1.32x better performance in read requests and up to 1.31x better write performance.



LearnedKV: Integrating LSM and Learned Index for Superior Performance on SSD



Thank you

2025.1.22

Presentation by Guangxun Zhao

GuangxunZhao@dankook.ac.kr