

CrossPrefetch: Accelerating I/O Prefetching for Modern Storage

S. Garg et al.

ASPLOS '24

2025. 02. 05

Presented by SeongHyeon Lee

leesh0812@dankook.ac.kr

Outline

- **Outline**
 - Introduction & Background
 - Motivation
 - Crossprefetch
 - Evaluation
 - Conclusion

Introduction & Background

- **Storage hardware is getting faster**
 - Newest NVMe SSDs have 7-10 GB/s throughput for sequential reads.
- **Applications have complex I/O access patterns**
 - Large storage bound apps transition between sequential, stride and random access pattern
- **Many filesystems specifically for modern storage hardware**
 - F2FS, noova, XFS etc.

Introduction & Background

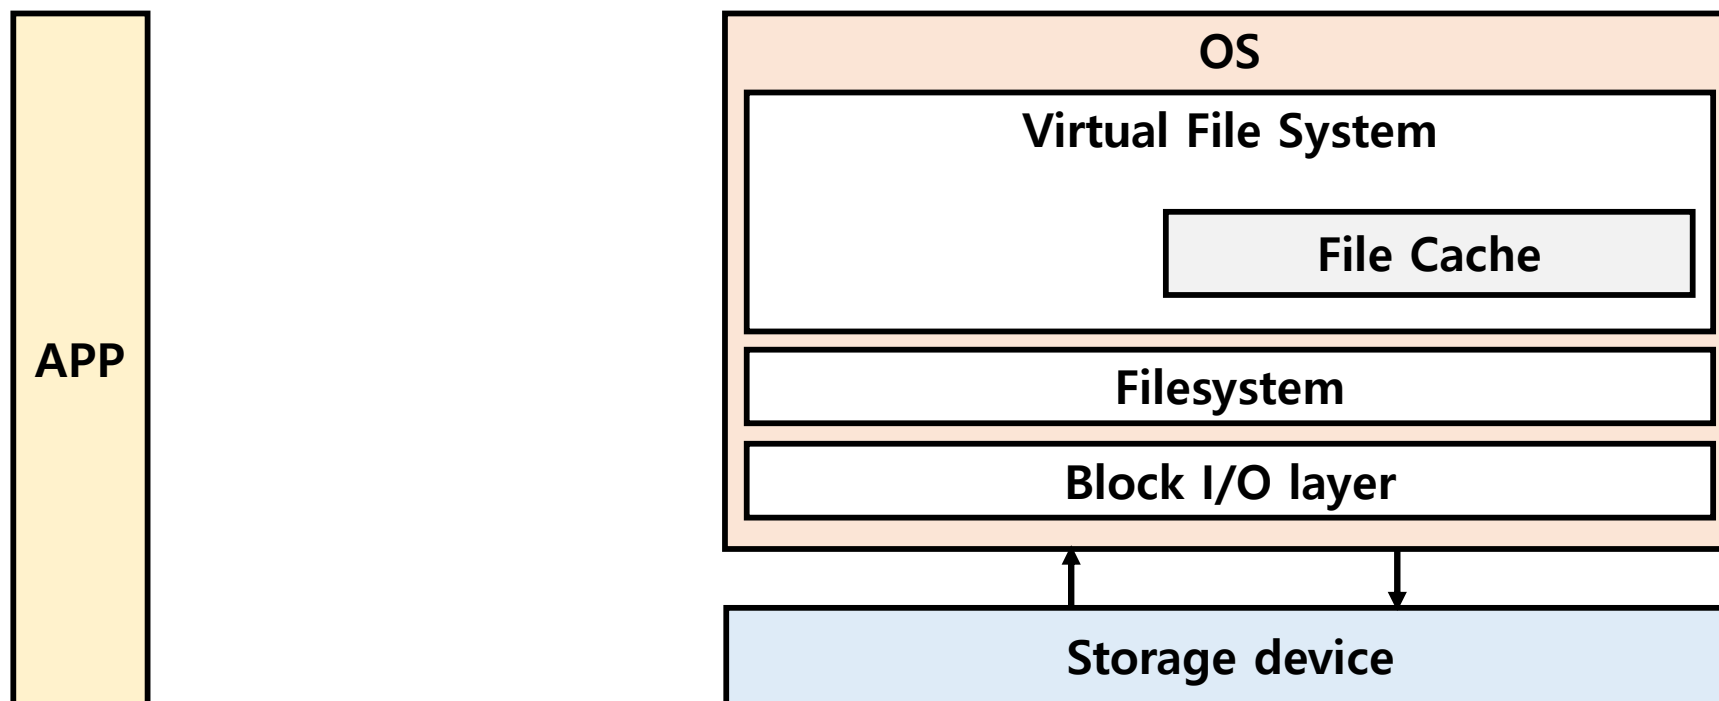


Readers Lock



Writers Lock

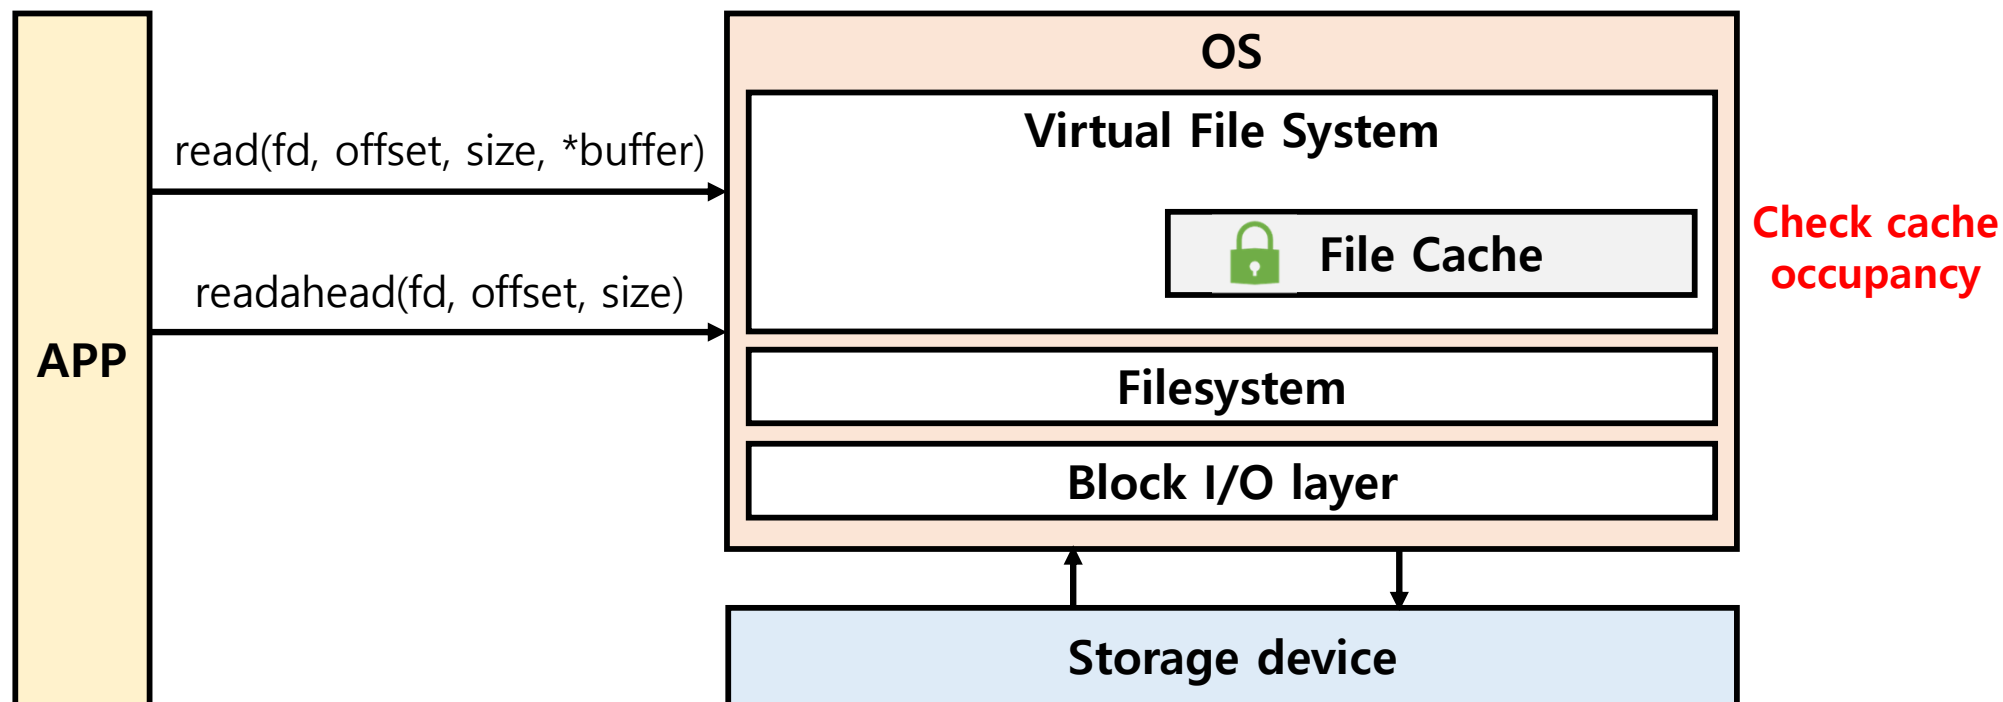
- Filesystem workflow



Introduction & Background

 Readers Lock
 Writers Lock

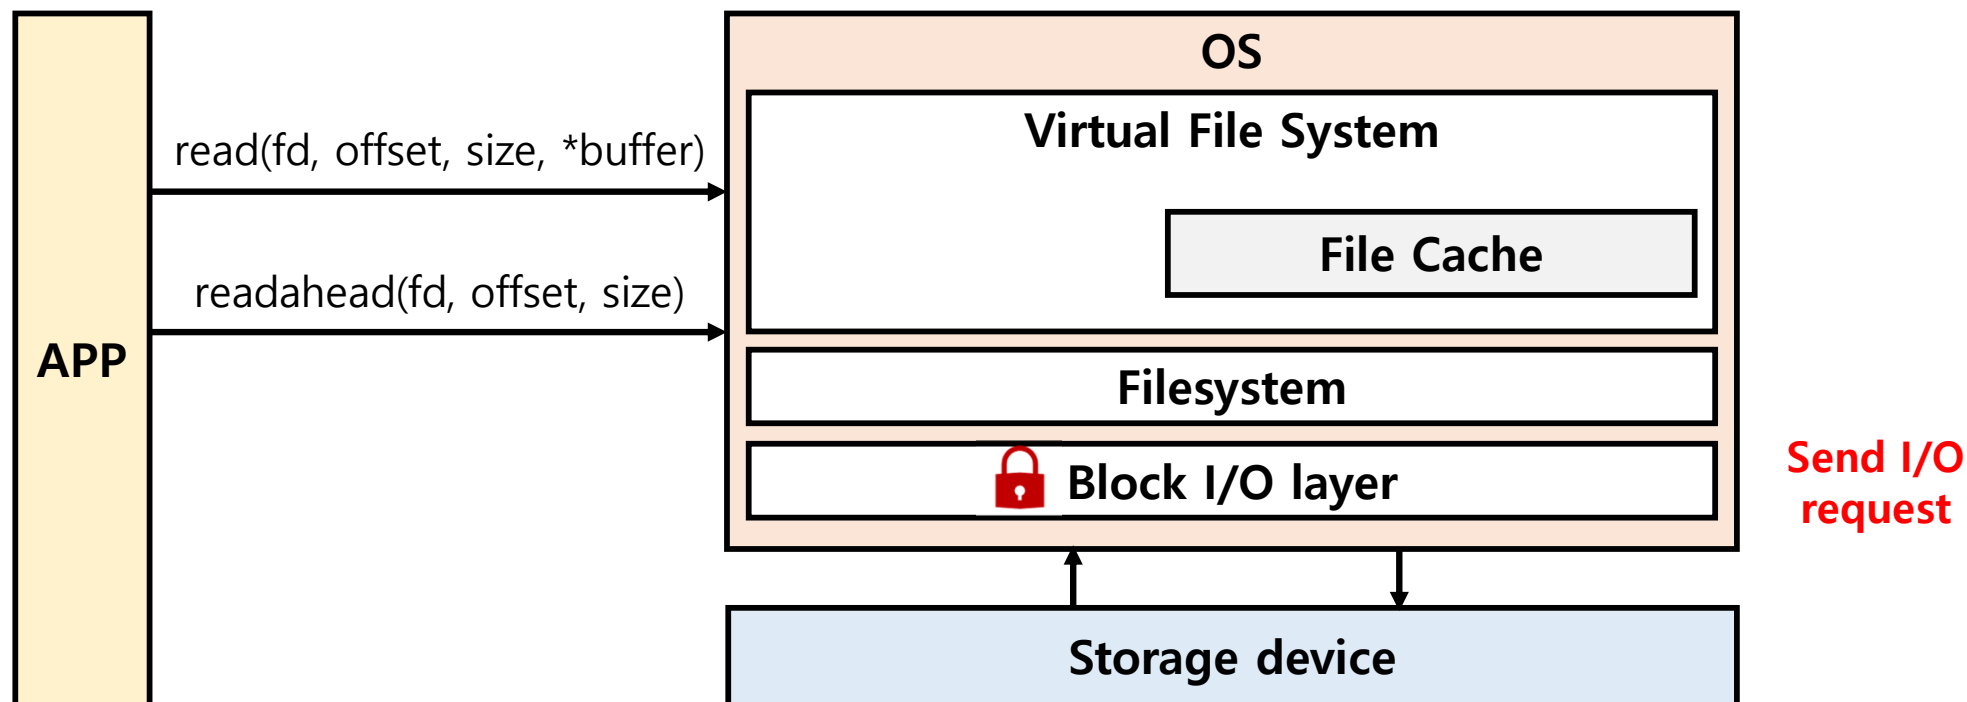
■ Filesystem workflow



Introduction & Background

 Readers Lock
 Writers Lock

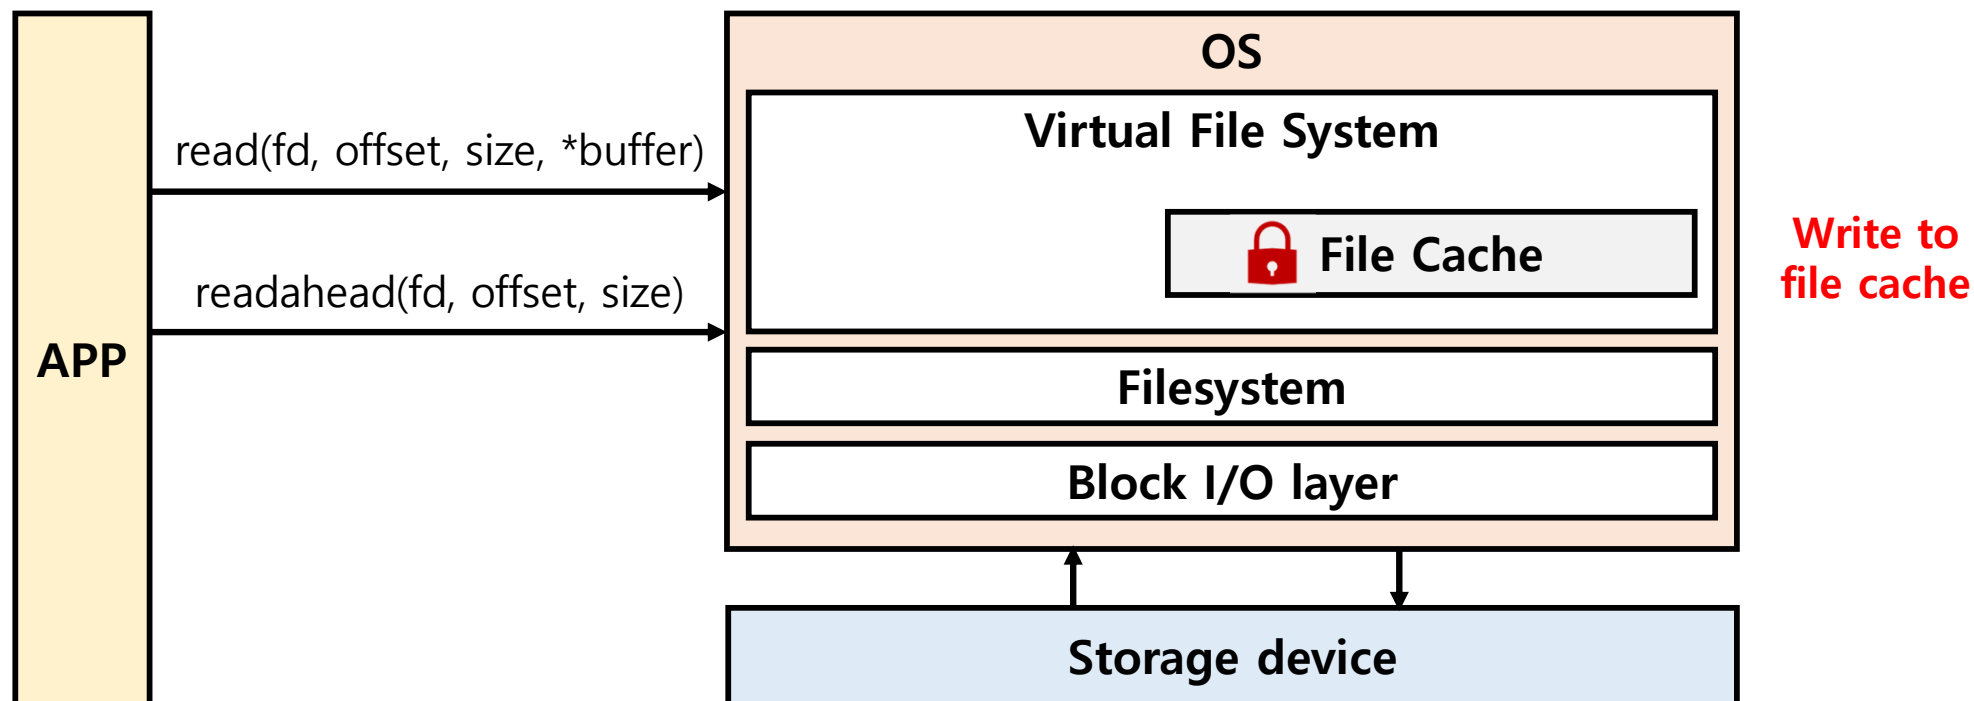
■ Filesystem workflow



Introduction & Background

 Readers Lock
 Writers Lock

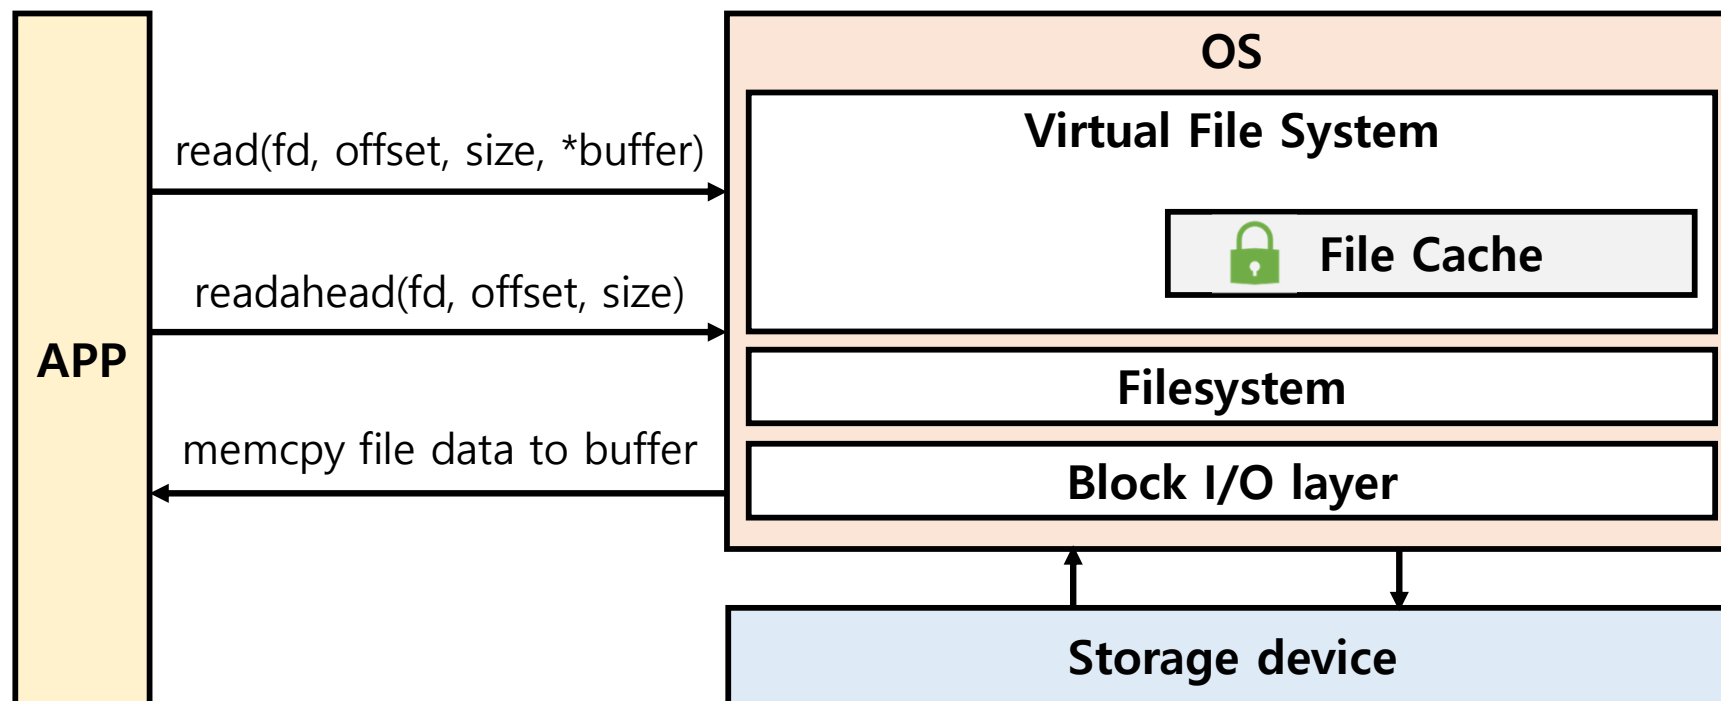
■ Filesystem workflow



Introduction & Background



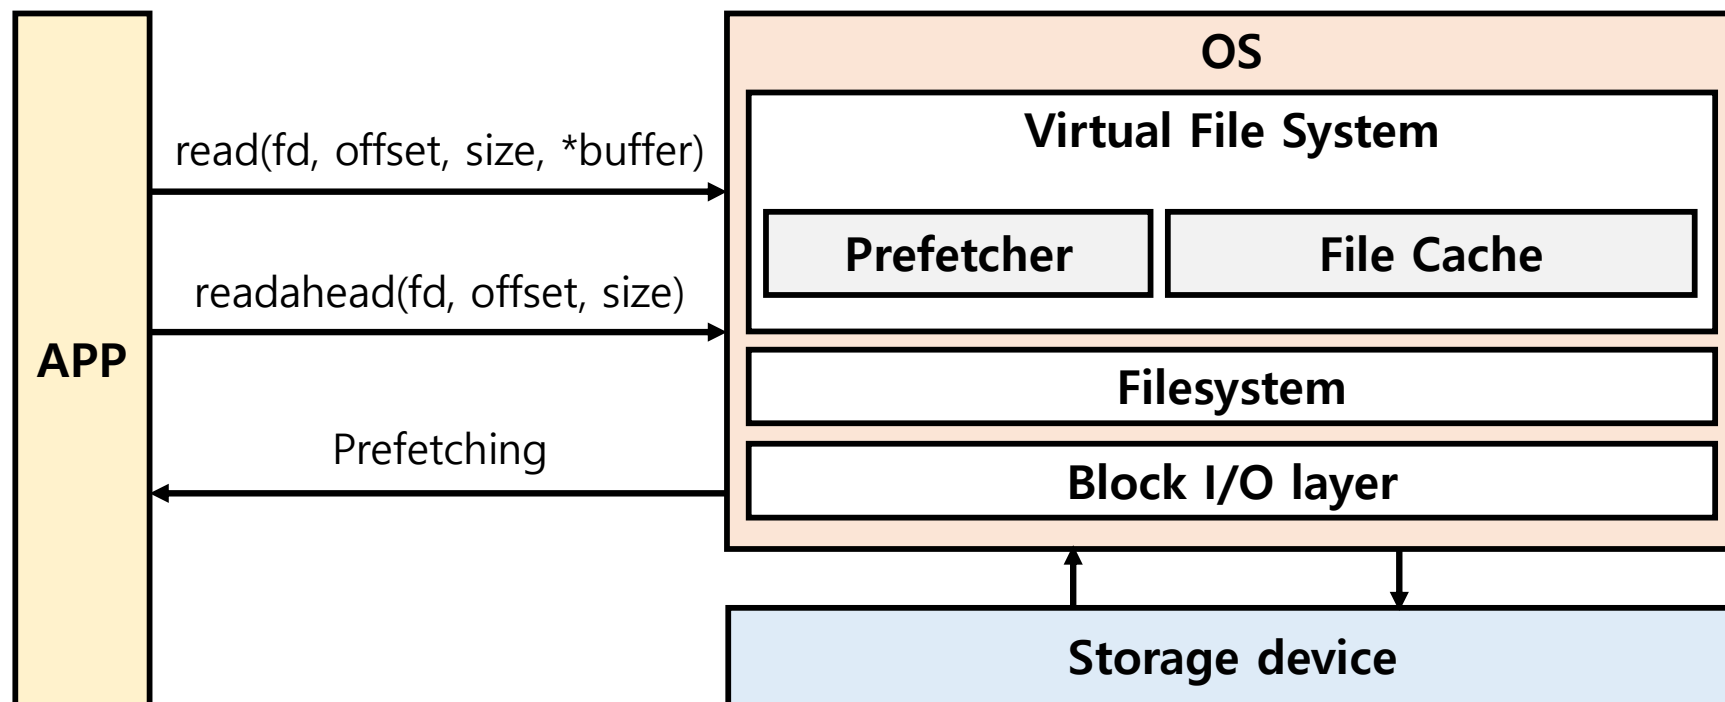
■ Filesystem workflow



Introduction & Background



■ Filesystem workflow



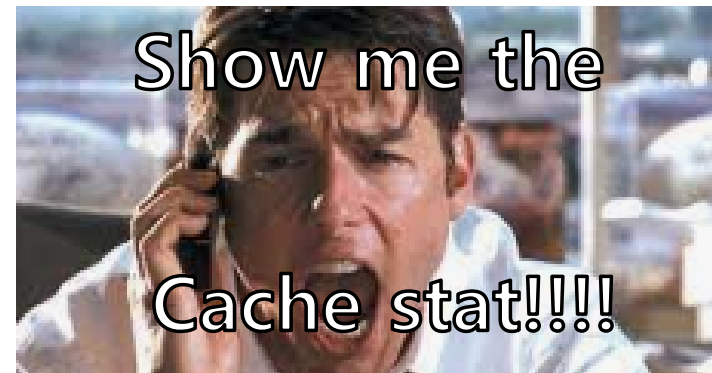
Introduction & Background

- **System call**

- int **fadvise**(int fd, off_t offset, off_t len, int advice)
- int **readahead**(int fd, off64_t offset, size_t count)

- **Advise**

- I/O access pattern : sequential, random
- Prefetching : will need, won't need



No information on what was prefetched to the application

→ Lack of cache visibility leads to poor prefetching and I/O performance

Introduction & Background



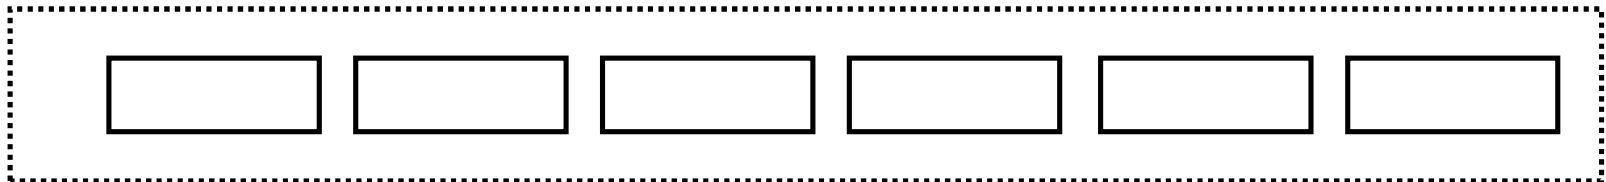
App requested for file page



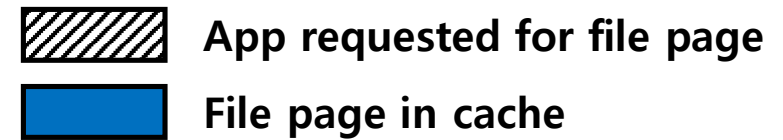
File page in cache

- Lack of I/O cache visibility : Under prefetching

Cache stat
for file



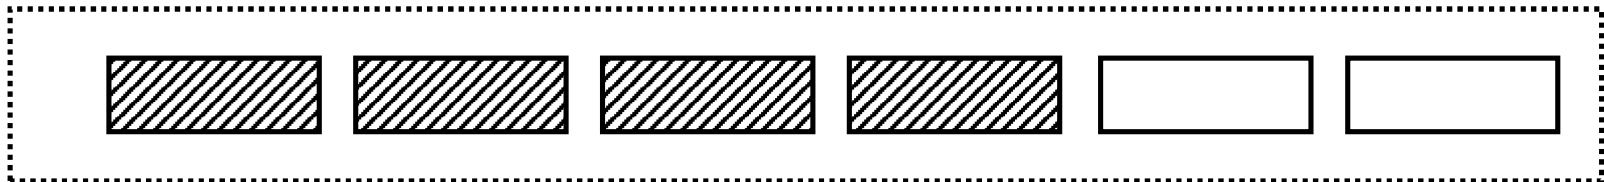
Introduction & Background



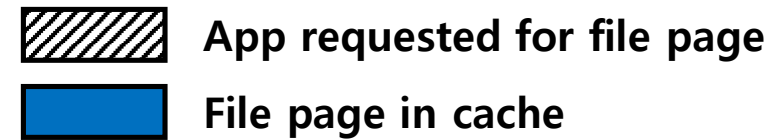
- Lack of I/O cache visibility : Under prefetching

Request 1: readahead(f1, start: pg1, sz: 4pg)

Cache stat
for file



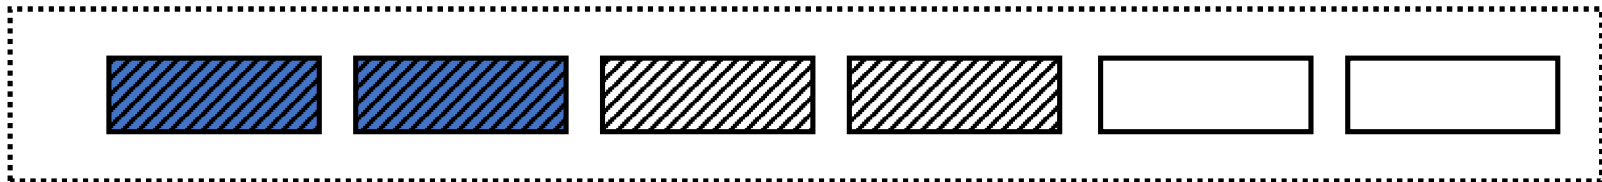
Introduction & Background



- Lack of I/O cache visibility : Under prefetching

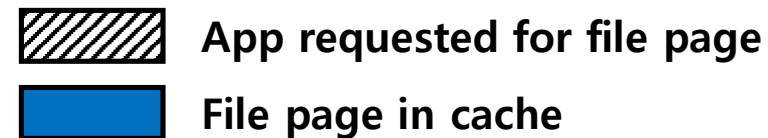
Request 1: readahead(f1, start: pg1, sz: 4pg)

Cache stat
for file



Request 1: Prefetched 2 pages

Introduction & Background

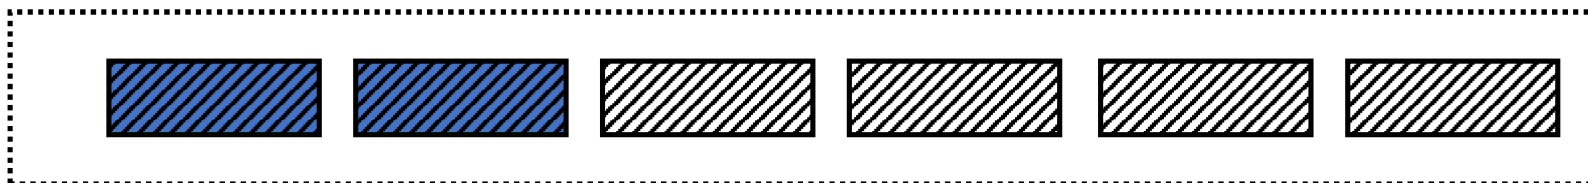


- Lack of I/O cache visibility : Under prefetching

Request 1: readahead(f1, start: pg1, sz: 4pg)

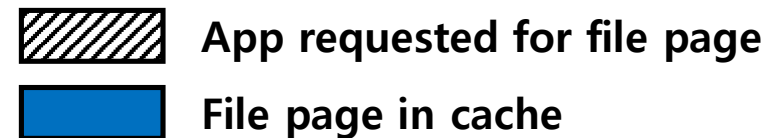
Request 2: readahead(f1, start: pg5, sz: 2pg)

Cache stat
for file



Request 1: Prefetched 2 pages

Introduction & Background

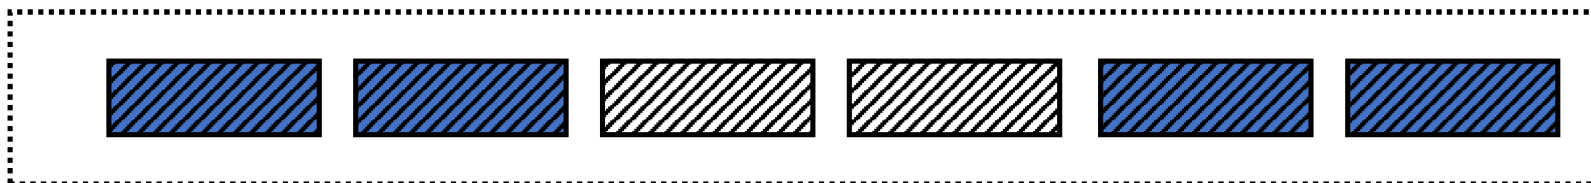


- Lack of I/O cache visibility : Under prefetching

Request 1: readahead(f1, start: pg1, sz: 4pg)

Request 2: readahead(f1, start: pg5, sz: 2pg)

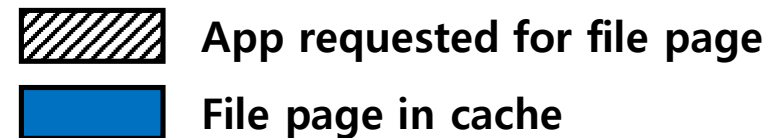
Cache stat
for file



Request 1: Prefetched 2 pages

Request 2: Prefetched 2 pages

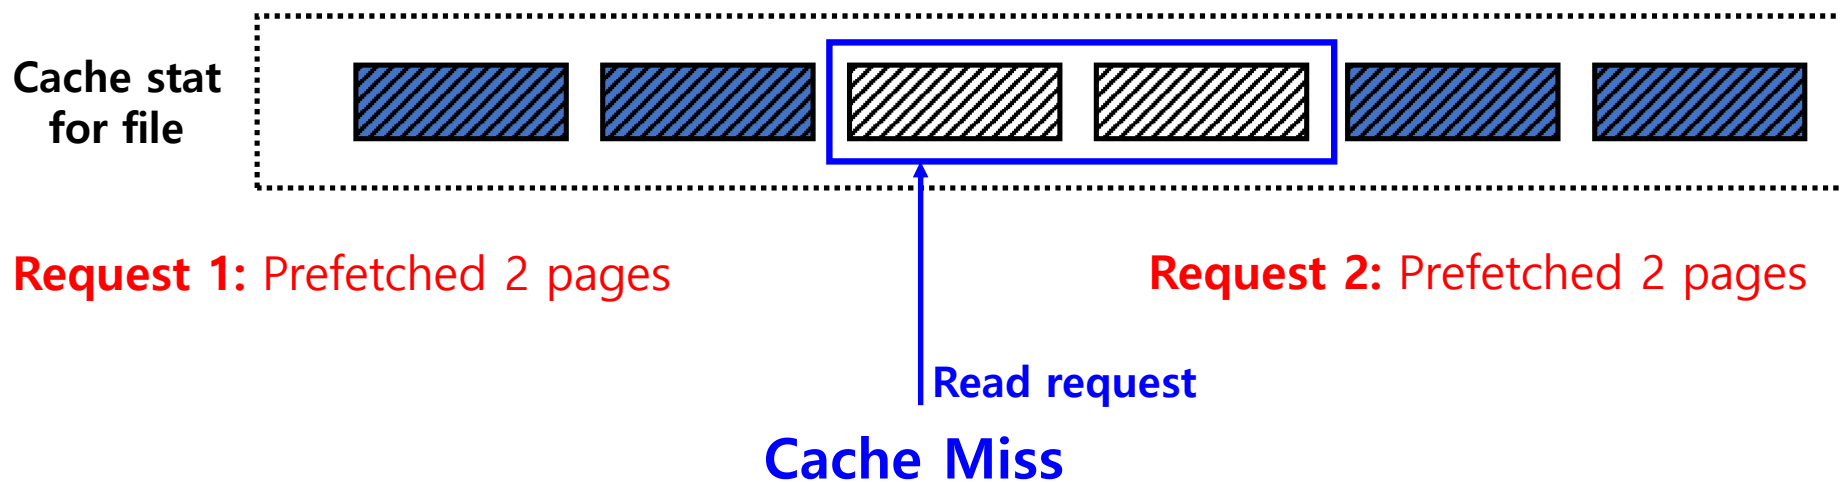
Introduction & Background



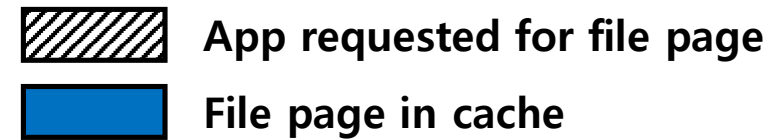
- Lack of I/O cache visibility : Under prefetching

Request 1: readahead(f1, start: pg1, sz: 4pg)

Request 2: readahead(f1, start: pg5, sz: 2pg)



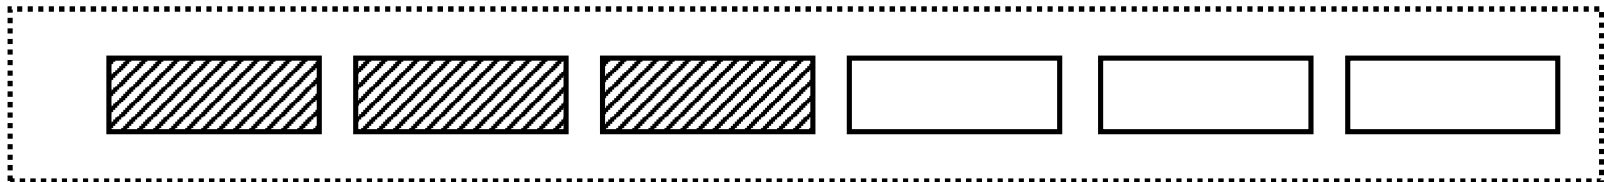
Introduction & Background



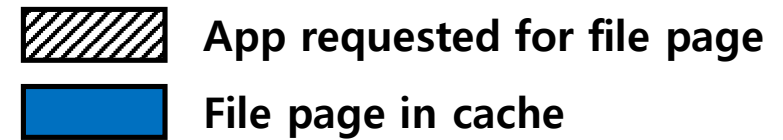
- Lack of I/O cache visibility : Over prefetching

Request 1: readahead(f1, start: pg1, sz: 3pg)

Cache stat
for file



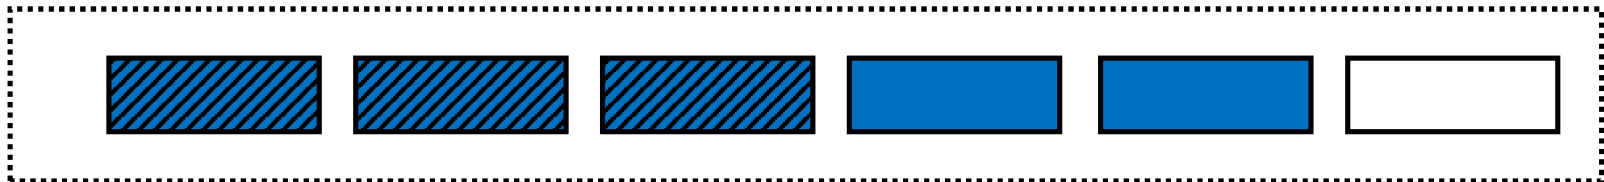
Introduction & Background



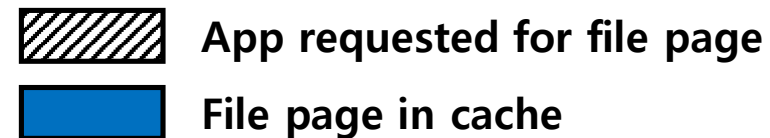
- Lack of I/O cache visibility : Over prefetching

Request 1: readahead(f1, start: pg1, sz: 3pg)

Cache stat
for file



Introduction & Background

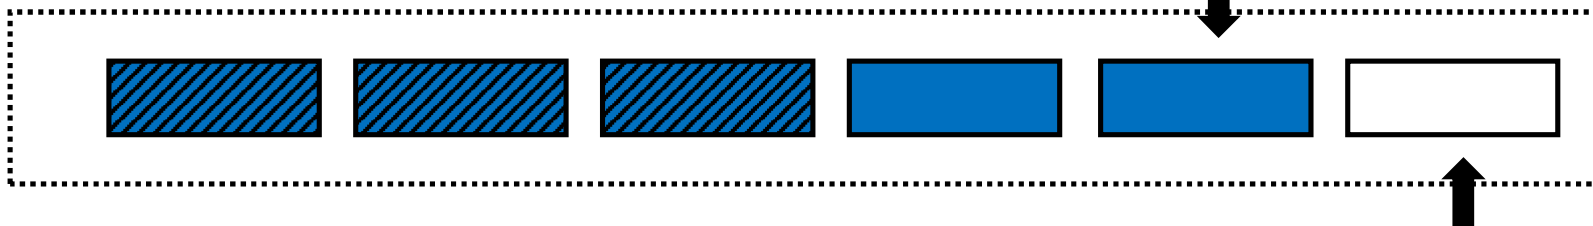


- Lack of I/O cache visibility : Over prefetching

Request 1: readahead(f1, start: pg1, sz: 3pg)

Request 2: readahead(f1, start: pg5, sz: 1pg)

Cache stat
for file



Read(f1, start: pg6, sz: 2KB)

Introduction & Background

▪ Related work

Prefetching Method	Description	Limitations
OS-based	Uses <code>readahead()</code> , <code>fadvise()</code> , and <code>madvise()</code> for prefetching	Applications cannot verify if prefetching was performed, rigid policies
Application-based	Prefetching implemented within applications like RocksDB, MySQL	Cannot check OS cache state, increases development complexity
Machine Learning (ML)-based	Uses Markov Chains, Reinforcement Learning (RL) for prefetching	Requires training time, lacks OS-application coordination
Compiler-based	Optimizes prefetching through static analysis	Struggles with dynamic workload changes, lacks OS cooperation

Motivation

- **Lack of Awareness of Prefetching Status**

- Applications cannot verify if prefetching was successful, leading to redundant or inefficient requests.

- **Concurrency Bottlenecks & Software Overhead**

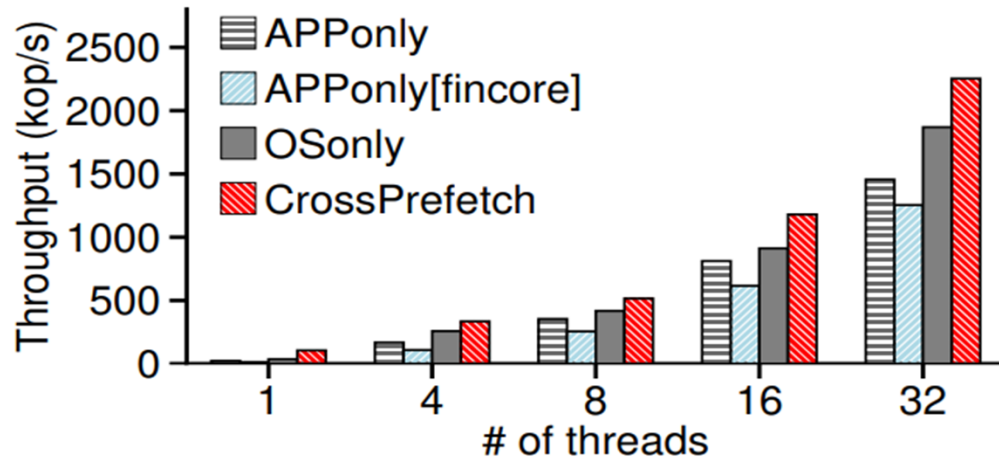
- Multiple threads sharing the same file can cause lock contention (e.g., Xarray lock in Linux), increasing system call overhead.

- **Inefficient Memory Utilization**

- Linux's incremental prefetching (128KB limit) does not adapt dynamically to available memory, leading to cache misses.

Motivation

▪ RocksDB Analysis



Oonly

where the OS handles prefetching

APPonly

disables OS and application prefetching for random access

APPonly[fincore]

uses a background prefetching thread to use fincore

	APPonly	APPonly[fincore]	OOnly	CrossPrefetch
Locking (%)	16	34	27	19
Cache Misses (%)	98.2	91.5	84.3	63.7

Table 1. Lock Overhead and Avg. Cache Misses (in %)

CrossPrefetch

- Disaggregate I/O prefetching responsibilities between the OS and a user-level runtime.
- Support concurrent prefetching and lightweight prediction.
- Enable aggressive prefetching and eviction without impacting memory budget.

CrossPrefetch

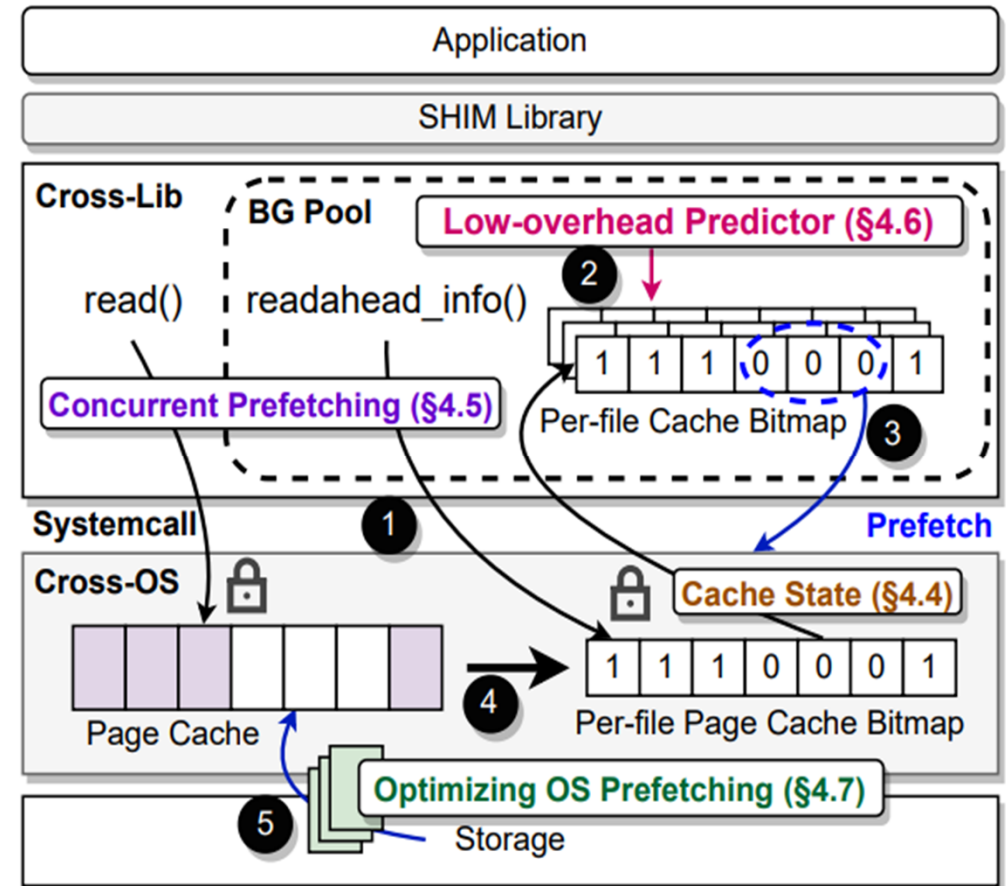
■ Cross-layered Approach

- Cross-OS

- Unlike traditional OS methods, manages **Per-file Page Cache Bitmap** at the inode level and optimizes prefetching based on this information.

- Cross-Lib

- Allows applications to directly control prefetching.



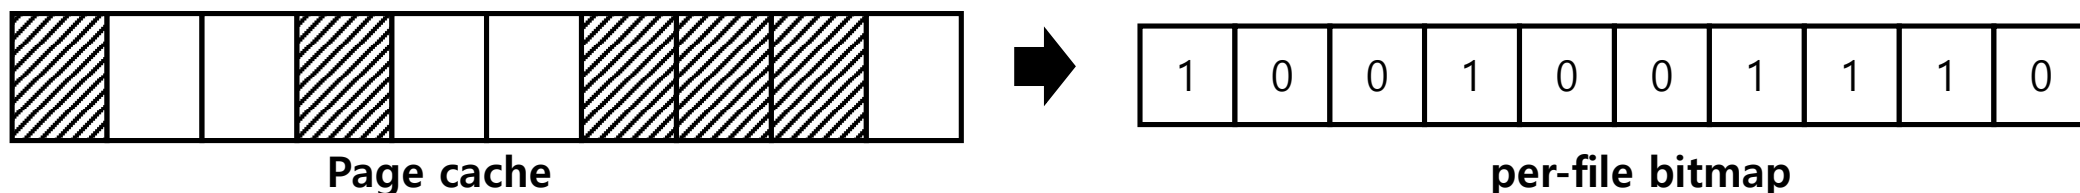
CrossPrefetch

- **Provide Visibility on I/O Prefetching State**
 - Overhead of scanning OS cache data structures
 - Cost of copying page cache state from OS to user space
 - Need for fast updates to per-file cache bitmap

CrossPrefetch

- **Provide Visibility on I/O Prefetching State**

- per-file bitmap



- readahead_info syscall

int **readahead**(int fd, off64_t offset, size_t count)

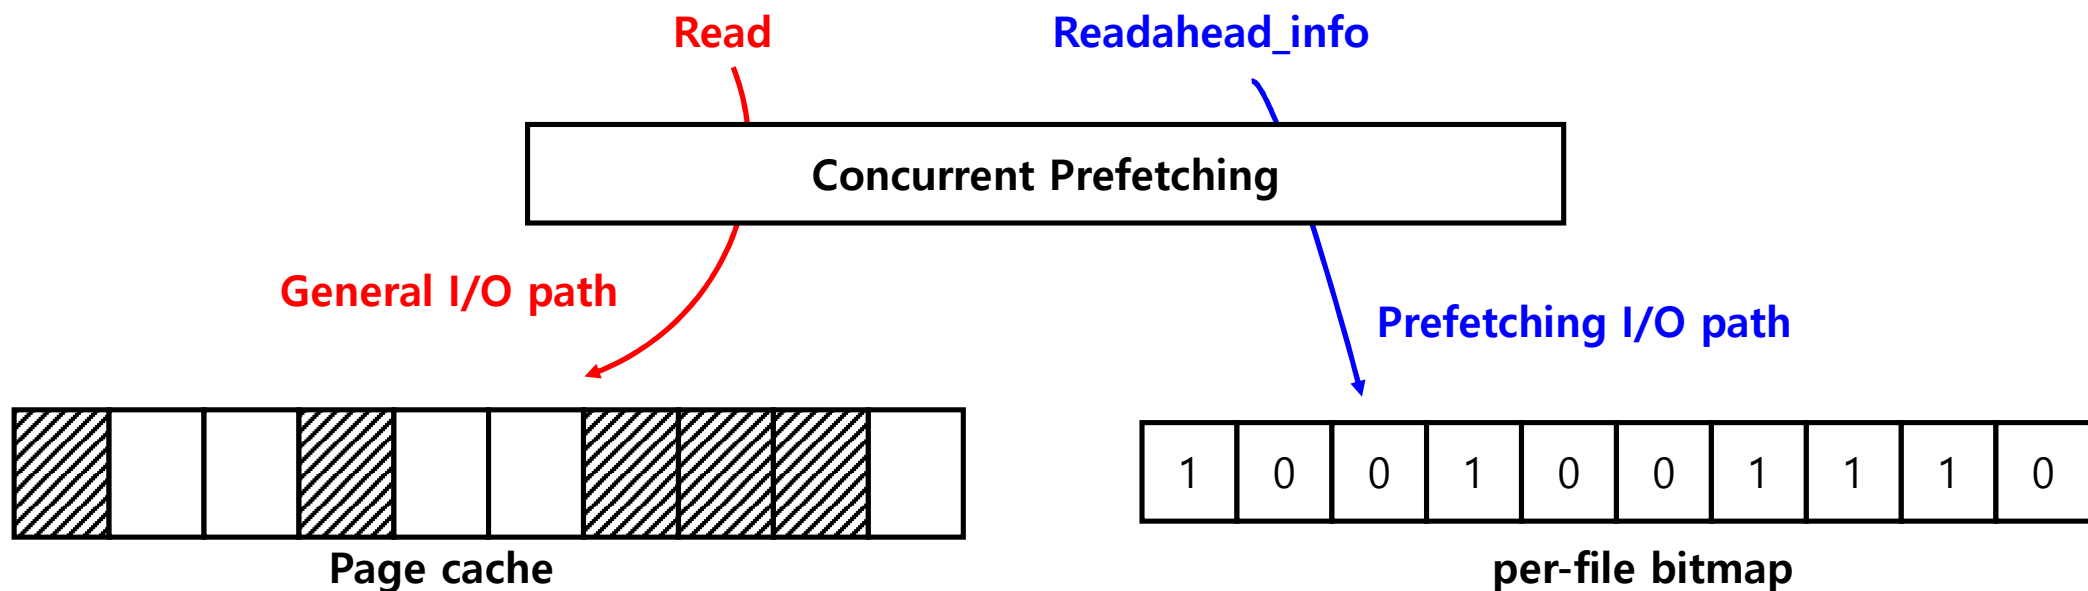


Int readahead_info(int fd, off64_t offset, size_t count, struct ra_info * info)

CrossPrefetch

■ Provide Visibility on I/O Prefetching State

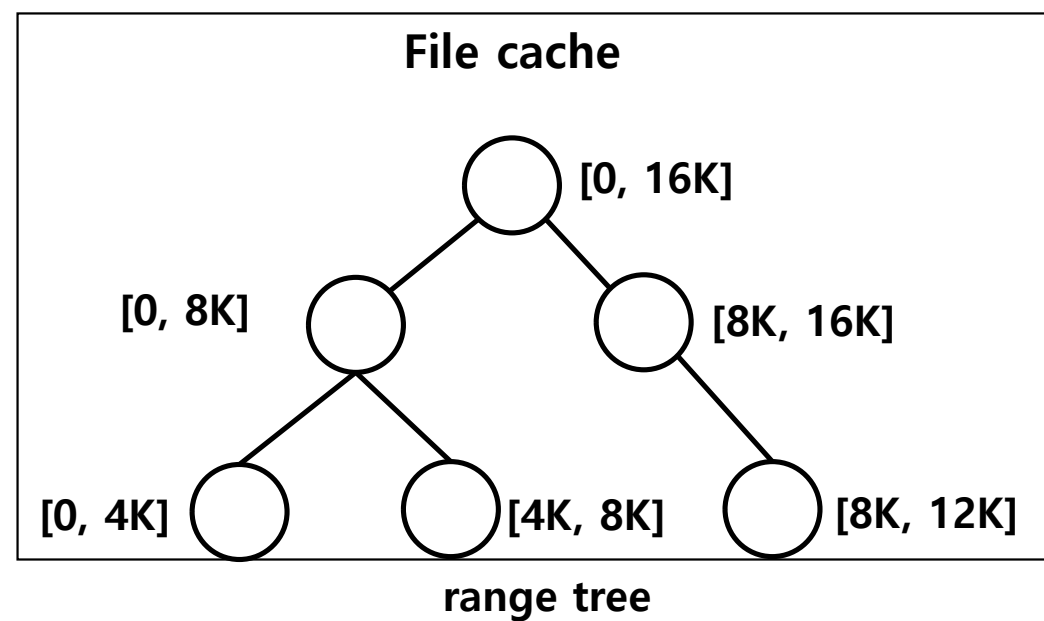
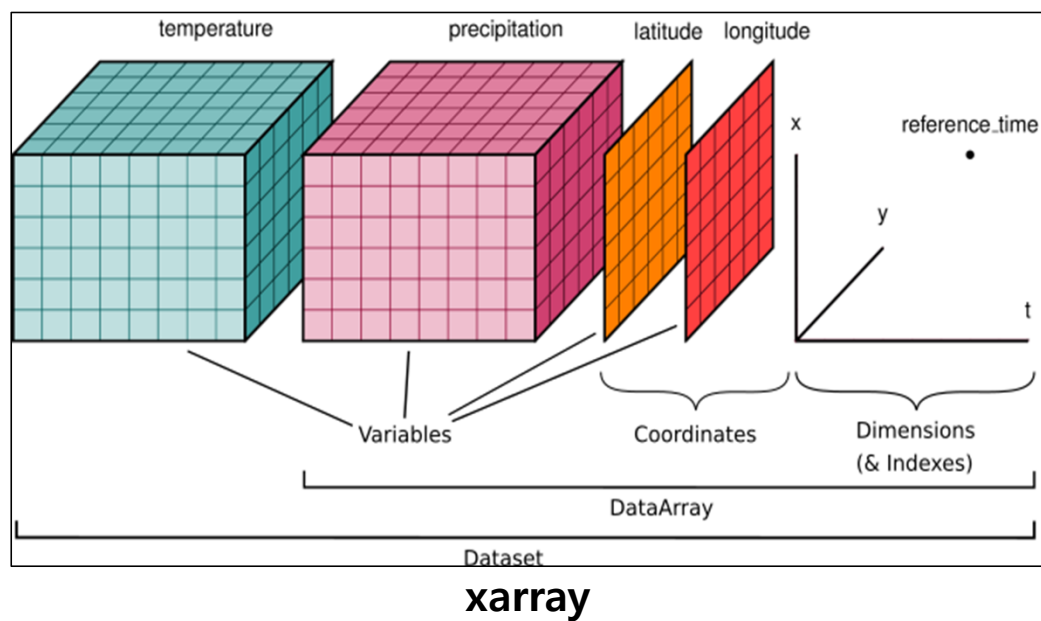
- Utilize per-file bitmaps to track cache status and improve prefetching efficiency
- Guarantee concurrency by separating the fast path and slow path



CrossPrefetch

- **Scalable and Concurrent Prefetching**

- Using per-file range tree



CrossPrefetch

Highly Random (000)	> 128KB distance
Random (001)	≤128KB distance
Partially Random(010)	Mixed sequential and random access
Likely Sequential (011)	Sequential access with occasional random accesses
Stride-based Sequential (100)	Sequential access with fixed strides
Definitely Sequential(110)	Continuous sequential access

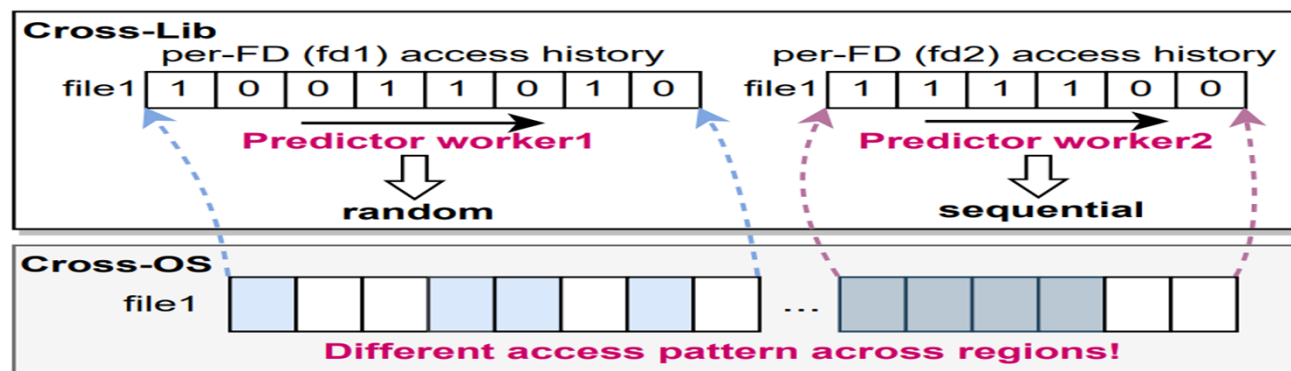
■ Low-overhead Prediction and Prefetching

- Pattern detector

- Uses a simple n-bit counter for detecting a file's access pattern

- Support for File-descriptor Prefetching

- Maintain an access pattern detector for each file descriptor
- The user-space file descriptor structure contains block range information and an access pattern counter



CrossPrefetch

Highly Random (000)	> 128KB distance	No prefetching
Random (001)	$\leq 128\text{KB}$ distance	128KB
Partially Random(010)	Mixed sequential and random access	512KB~1MB
Likely Sequential (011)	Sequential access with occasional random accesses	2MB
Stride-based Sequential (100)	Sequential access with fixed strides	512KB~1MB
Definitely Sequential(110)	Continuous sequential access	2MB

■ Low-overhead Prediction and Prefetching

- Memory-aware Aggressive Prefetching and Eviction

- Perform aggressive prefetching from the beginning of execution to reduce initial cache misses
- Adjust prefetching aggressiveness by setting a High Threshold and a Low Threshold
- Assume sequential access when a file is first opened

- Aggressive Reclamation

- Reclaim inactive file pages based on the LRU policy
- Optimize cache usage for large files

- Support for Memory-Mapped I/O

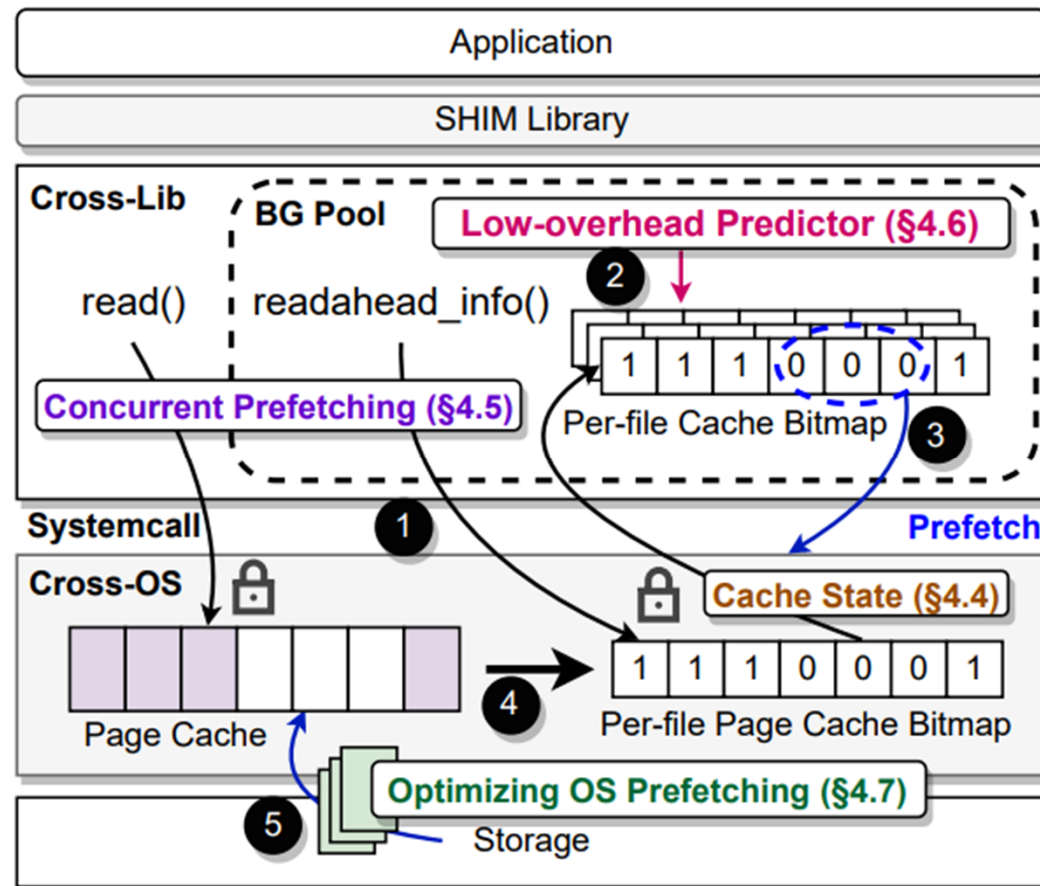
- Utilize the OS-provided cache bitmap to detect access patterns
- Analyze access patterns using background threads and apply an appropriate prefetching window size
- However, this approach is similar to existing Linux OS prefetching and may have lower accuracy

CrossPrefetch

▪ Optimizing OS Prefetching Path

- Linux limits the incremental prefetch threshold to 32 pages (128KB)
- Extend the OS to dynamically increase the prefetch threshold using the info structure in the readahead_info system call
 - Restrict actual prefetch request sizes to a maximum of 64MB
 - ✓ Excessively large prefetch requests may degrade blocking I/O (read/write) performance

CrossPrefetch



Evaluation

- **Experimental Environment**

- **Configuration**

- CPU: Intel Xeon Gold 6226R (16 cores, 2.9GHz)
 - Memory: 128GB DDR4
 - Storage: Samsung PM1733 NVMe SSD (6.4TB)

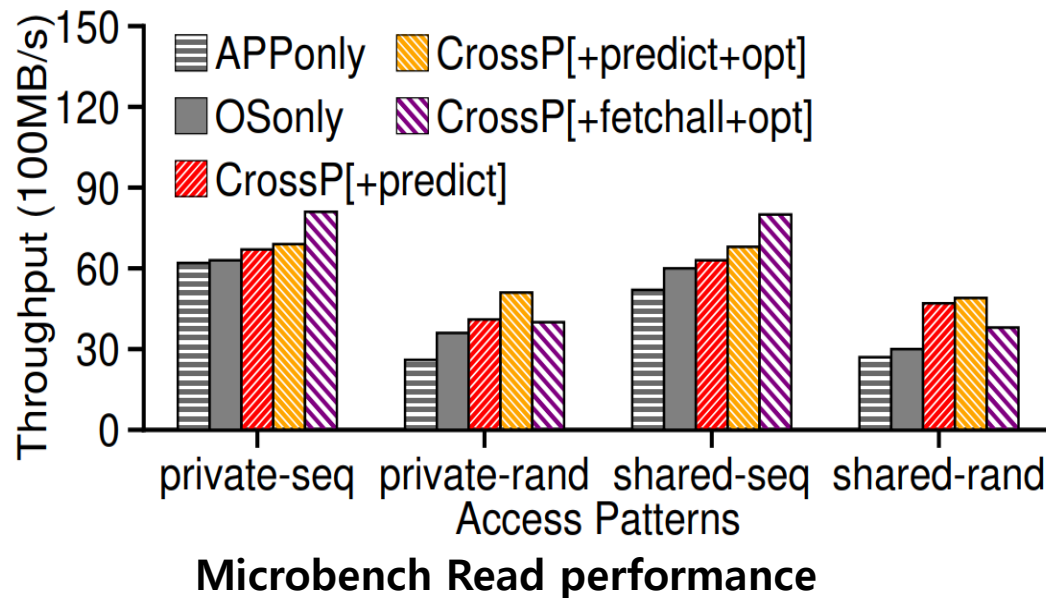
- **Benchmark Tools**

- fio, RocksDB, TPC-C (OLTP workloads)

Evaluation

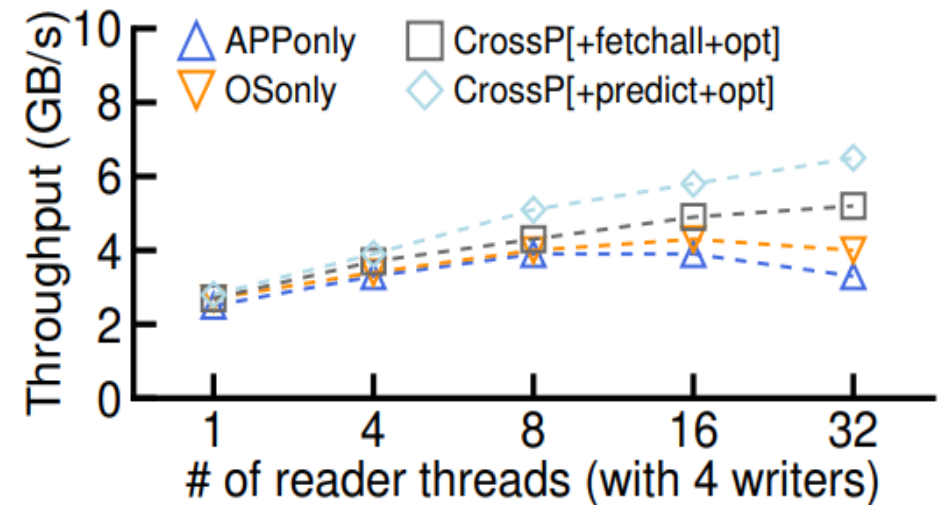
■ Microbench

- sequential & random access



Mechanism	Description
<i>APPonly</i>	Application tailored prefetching using readahead calls
<i>OSonly</i>	Prefetching delegated to OS and application prefetching is disabled
<i>CrossP[+predict]</i>	Fine-grained and low-interference prediction avoiding prefetching entire file
<i>CrossP[+predict+opt]</i>	<i>CrossP[+predict]</i> without OS limits but also provide memory-centric aggressive prefetching and eviction
<i>CrossP[+fetchall+opt]</i> (memory insensitive)	Proposed CrossPrefetch that uses cache state awareness to prefetch missing blocks of a file using <code>readahead_info()</code> assumes all data fits in memory

Table 2. Comparison Approaches

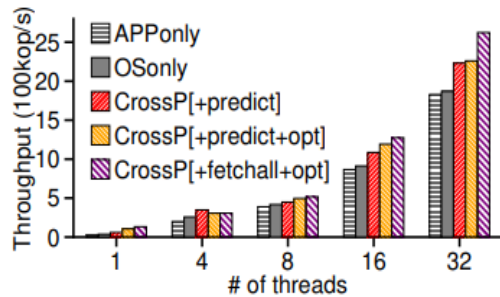


Evaluation

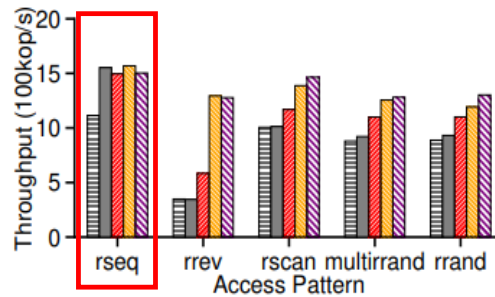
- **Local NVMe**
 - RocksDB dbbench

Mechanism	Description
<i>APPonly</i>	Application tailored prefetching using readahead calls
<i>OSonly</i>	Prefetching delegated to OS and application prefetching is disabled
<i>CrossP[+predict]</i>	Fine-grained and low-interference prediction avoiding prefetching entire file
<i>CrossP[+predict+opt]</i>	<i>CrossP[+predict]</i> without OS limits but also provide memory-centric aggressive prefetching and eviction
<i>CrossP[+fetchall+opt]</i> (memory insensitive)	Proposed CrossPrefetch that uses cache state awareness to prefetch missing blocks of a file using <code>readahead_info()</code> assumes all data fits in memory

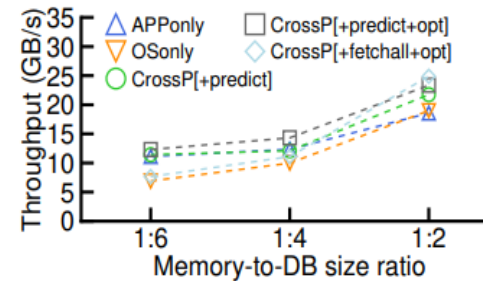
Table 2. Comparison Approaches



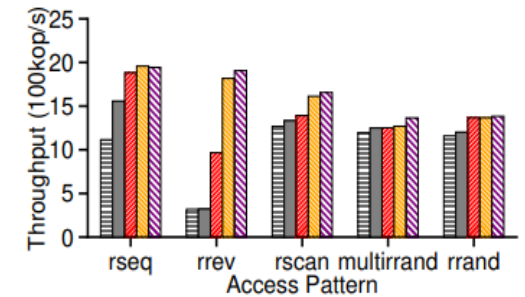
(a) Multi-read Random on ext4



(b) Access Patterns on ext4



(c) Memory capacity impact



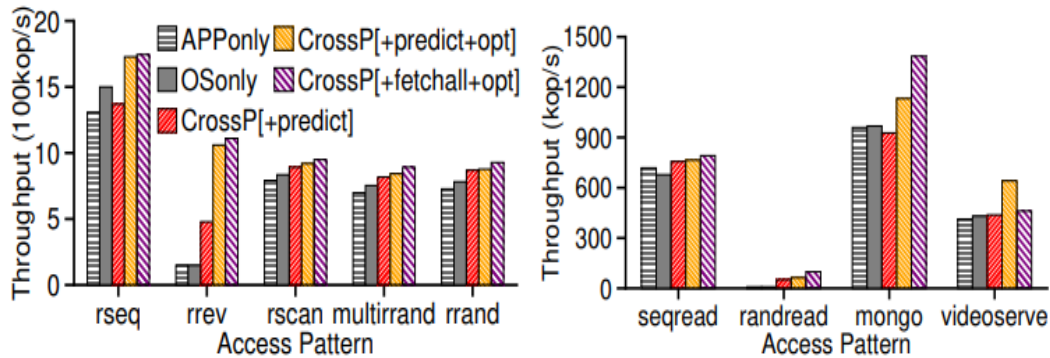
(d) Access Patterns on F2FS

RocksDB DBbench on Local NVMe

Evaluation

Local & Remote NVMe

- RocksDB dbbench, Filebench, YCSB, Snappy



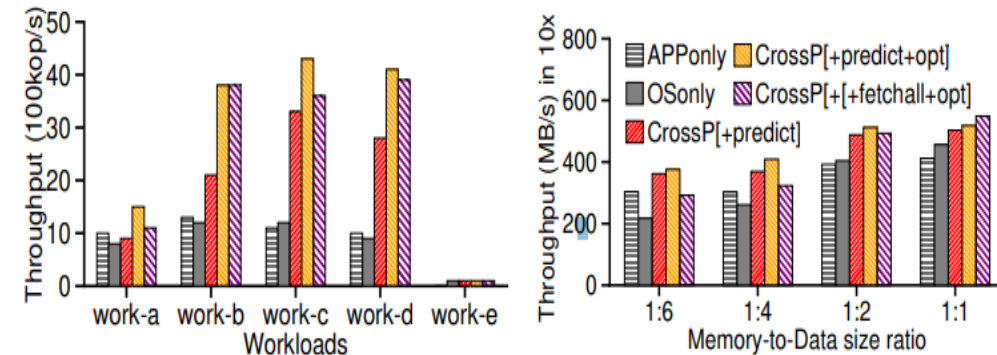
(a) Access Patterns

(b) Filebench Workloads

RocksDB on Remote NVMe (8a) and Filebench (8b)

Mechanism	Description
<i>APPonly</i>	Application tailored prefetching using readahead calls
<i>OSonly</i>	Prefetching delegated to OS and application prefetching is disabled
<i>CrossP[+predict]</i>	Fine-grained and low-interference prediction avoiding prefetching entire file
<i>CrossP[+predict+opt]</i>	<i>CrossP[+predict]</i> without OS limits but also provide memory-centric aggressive prefetching and eviction
<i>CrossP[+fetchall+opt]</i> (memory insensitive)	Proposed CrossPrefetch that uses cache state awareness to prefetch missing blocks of a file using readahead_info() assumes all data fits in memory

Table 2. Comparison Approaches



(a) YCSB workload

(b) Snappy

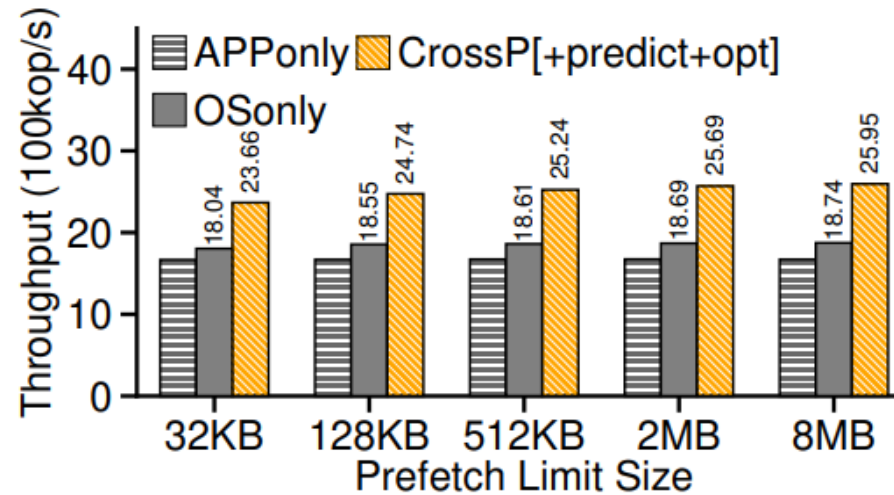
Real-world Workloads

Evaluation

■ Prefetch Limit Impact

Mechanism	Description
<i>APPonly</i>	Application tailored prefetching using readahead calls
<i>OSonly</i>	Prefetching delegated to OS and application prefetching is disabled
<i>CrossP[+predict]</i>	Fine-grained and low-interference prediction avoiding prefetching entire file
<i>CrossP[+predict+opt]</i>	<i>CrossP[+predict]</i> without OS limits but also provide memory-centric aggressive prefetching and eviction
<i>CrossP[+fetchall+opt]</i> (memory insensitive)	Proposed CrossPrefetch that uses cache state awareness to prefetch missing blocks of a file using <code>readahead_info()</code> assumes all data fits in memory

Table 2. Comparison Approaches



Prefetch Limit Impact

Conclusion

- **Lack of communication between the OS and userspace**
 - The lack of direct communication between the OS and userspace in traditional prefetching mechanisms results in **inefficient cache management and unnecessary I/O operations**.
 - **Conflicting motivations and assumptions** between the OS and applications reduce overall system performance.
- **Harmony between layers in the system reduces friction**
 - **CrossPrefetch** enhances coordination between OS-level and userspace prefetching by **sharing per-file cache state and prefetching hints**.
 - Sharing "Need-To-Know" information between layers significantly improves prefetch accuracy and system performance.

Q&A

Q&A